

Global AI Accountability Protocol (GAAP) 2.2

Comprehensive Ethical Framework for Responsible AI Development and Governance

Release Date: July 2025

Version: 2.2

Status: Active Standard

Table of Contents

1. Executive Summary
2. Introduction
3. Background and Evolution
4. Core Principles
5. Implementation Tiers & Requirements
6. Enforcement Mechanisms
7. Certification Process
8. Appendices
9. Glossary & Definitions
10. Contact & Oversight
11. References

Executive Summary

The Global AI Accountability Protocol (GAAP) 2.2 represents a milestone in establishing universal standards for ethical AI development and deployment. As artificial intelligence systems become increasingly integrated into critical decision-making processes across healthcare, finance, governance, and social infrastructure, the need for a comprehensive, globally applicable framework has never been more urgent.

GAAP 2.2 addresses the evolving landscape of AI ethics through a secular, evidence-based approach that transcends cultural, religious, and philosophical boundaries. This framework introduces groundbreaking provisions for neuro-rights protection, environmental sustainability, enhanced whistleblower mechanisms, and cross-jurisdictional interoperability—responding directly to the challenges emerging in 2025's AI ecosystem.

Key Innovations in Version 2.2

- **Neuro-Rights Framework:** First comprehensive protocol for brain-computer interface (BCI) data governance and cognitive liberty protection
- **Environmental Accountability:** Mandatory carbon impact disclosure and energy-efficient operation requirements
- **Enhanced Whistleblower Protections:** 72-hour mandatory escalation with automatic external notification mechanisms
- **Global Interadoptability:** Cross-certification with GDPR, CPRA, LGPD, PIPEDA, and international standards
- **Contextual Harm Assessment:** Regional support accessibility adjustments for harm severity categorization
- **Indigenous Data Sovereignty:** Explicit requirements for OCAP® compliance and consultation frameworks

Who Should Adopt GAAP 2.2

This protocol is designed for organizations across all sectors deploying AI systems, including:

- Technology companies developing AI products and services
- Healthcare institutions using AI for diagnosis, treatment, or patient management
- Financial institutions deploying AI for risk assessment, fraud detection, or customer service
- Government agencies implementing AI in public services or critical infrastructure
- Research institutions advancing AI capabilities
- Military and defense organizations utilizing AI systems

GAAP 2.2 employs a tiered implementation structure that scales requirements based on system risk, scope, and impact—ensuring proportionate governance without stifling innovation.

1. Introduction

1.1 The Need for Global AI Accountability

Artificial intelligence has transitioned from experimental technology to foundational infrastructure underpinning modern society. From algorithmic content curation affecting billions of users to AI-assisted medical diagnostics determining treatment pathways, these systems wield unprecedented influence over human welfare, rights, and opportunities.

However, this rapid proliferation has occurred largely without unified ethical guardrails. The result has been a patchwork of regional regulations, corporate self-governance initiatives, and reactive policy responses to AI-related harms. High-profile incidents—ranging from biased hiring algorithms to facial recognition misidentifications leading to wrongful arrests—have demonstrated the urgent need for proactive, comprehensive standards.

1.2 The GAAP Approach: Secular and Universal

The term "secular" in GAAP's context emphasizes **universal applicability regardless of cultural, religious, or philosophical backgrounds**. Unlike frameworks rooted in specific ethical traditions or regional values, GAAP 2.2 establishes baseline requirements grounded in empirically demonstrable principles of harm reduction, transparency, fairness, and human dignity.

This approach ensures that GAAP can be adopted globally without exclusion or bias, providing a common language for AI governance across diverse jurisdictions while respecting legitimate regional variations in implementation.

1.3 Multi-Tier Stress Testing Methodology

GAAP 2.2 has been developed through rigorous multi-tier stress testing involving:

- **Adversarial Red Teaming:** Simulated attacks on AI systems to identify vulnerabilities in safety mechanisms, bias mitigation, and data protection
- **Cross-Cultural Validation:** Testing protocol effectiveness across diverse demographic groups, cultural contexts, and regulatory environments
- **Scenario-Based Analysis:** Evaluation of framework application to edge cases, emerging technologies, and hypothetical crisis situations
- **Stakeholder Consultation:** Input from AI developers, ethicists, policymakers, affected communities, industry leaders, and civil society organizations
- **Technical Feasibility Assessment:** Verification that all requirements are technically implementable with current or near-term capabilities

This iterative development process ensures GAAP 2.2 addresses both current challenges and anticipated future developments in AI technology.

1.4 Scope and Limitations

GAAP 2.2 applies to AI systems across all domains and use cases, with tiered requirements scaling to system risk and impact. The protocol covers:

- Machine learning models (supervised, unsupervised, reinforcement learning)
- Natural language processing systems
- Computer vision and multimodal systems
- Robotic systems with AI components
- Decision support and autonomous decision-making systems
- Brain-computer interfaces with AI integration

Out of Scope: GAAP 2.2 does not regulate non-AI algorithmic systems without learning components, purely statistical models, or human-only decision-making processes. However, hybrid systems combining traditional algorithms with AI components fall under GAAP jurisdiction.

2. Background and Evolution

2.1 The Path to GAAP 2.2

The Global AI Accountability Protocol emerged from recognition that fragmented, jurisdiction-specific regulations were insufficient for governing inherently borderless AI technologies. Early versions of GAAP (1.0 in 2022, 2.0 in 2024) established foundational principles of transparency, fairness, and accountability that gained widespread adoption across industry and government.

Key Milestones:

- **GAAP 1.0 (2022):** Established core principles and basic transparency requirements
- **GAAP 2.0 (2024):** Introduced tiered implementation structure and bias mitigation protocols
- **GAAP 2.1 (Early 2025):** Added preliminary environmental considerations and cross-border data governance
- **GAAP 2.2 (July 2025):** Comprehensive neuro-rights framework, enhanced whistleblower protections, and contextual harm assessment

2.2 Driving Forces Behind Version 2.2

Several developments in late 2024 and early 2025 necessitated the comprehensive update represented by GAAP 2.2:

Neuro-Technology Proliferation: Commercial brain-computer interfaces transitioned from research novelty to consumer products, creating urgent need for cognitive liberty protections and neural data governance frameworks.

Environmental Impact Awareness: Recognition of AI's carbon footprint—with large language model training consuming energy equivalent to hundreds of households annually—demanded mandatory sustainability disclosures and efficiency requirements.

Whistleblower Crisis: High-profile cases of internal AI safety concerns being suppressed or ignored highlighted need for stronger protection mechanisms and mandatory external escalation pathways.

Cross-Border AI Deployment: Increasing global deployment of AI systems created friction between regional data protection regimes, requiring harmonized cross-certification standards.

Contextual Harm Realization: Recognition that identical AI failures produce vastly different real-world impacts depending on regional infrastructure and support availability, necessitating context-sensitive harm assessment.

2.3 Relationship to Existing Frameworks

GAAP 2.2 is designed to complement and interoperate with existing regulatory frameworks rather than replace them:

Framework	Region	GAAP 2.2 Relationship
GDPR (General Data Protection Regulation)	European Union	GAAP data governance requirements align with and extend GDPR principles; cross-certification available
CPRA (California Privacy Rights Act)	California, USA	GAAP Tier 2+ requirements encompass CPRA obligations; mutual recognition agreements
LGPD (Lei Geral de Proteção de Dados)	Brazil	GAAP data sovereignty principles compatible with LGPD; unified audit criteria
PIPEDA (Personal Information Protection and Electronic Documents Act)	Canada	GAAP consent mechanisms satisfy PIPEDA requirements; cross-recognition pathways
ISO/IEC 23894	International	GAAP certification includes ISO/IEC 23894 risk management verification
IEEE 7000	International	GAAP ethical considerations incorporate IEEE 7000 systems design principles

Organizations certified under GAAP 2.2 receive streamlined compliance pathways for these complementary frameworks, reducing duplication and administrative burden.

3. Core Principles

GAAP 2.2 is founded on ten core principles that serve as both philosophical foundation and practical implementation guide. These principles are not hierarchical—all must be satisfied for full compliance—but are numbered for organizational clarity.

3.1 Transparency & Explainability

Principle Statement: AI systems must produce auditable outputs with clear, user-accessible explanations appropriate to audience technical sophistication.

Rationale: Transparency serves multiple critical functions: enabling accountability when systems cause harm, facilitating informed consent from affected individuals, supporting regulatory oversight, and building public trust in AI deployment.

Implementation Requirements:

User-Accessible Explanations

- **Plain-Language Summaries:** Non-technical explanations available to all end-users in their primary language, describing what the AI does, what data it uses, and how decisions are reached
- **Example:** A credit scoring AI must explain "Your score was lowered because of high credit utilization (using 80% of available credit) and recent late payments" rather than providing only a numeric score

Technical Documentation

- **Detailed Technical Specifications:** Comprehensive documentation of model architecture, training procedures, data sources, and performance metrics available to auditors and regulators
- **Optional Technical Dashboards:** Interactive tools allowing technical users to explore model behavior, feature importance, and decision boundaries

Tier-Specific Requirements

- **Tier 1:** Basic output logging and plain-language explanations
- **Tier 2+:** External-facing explainability tools allowing affected individuals to request and receive specific explanations for decisions impacting them
- **Tier 3+:** Real-time explainability APIs for third-party audit tools
- **Tier 4:** Complete model interpretability with formal verification of explanation accuracy

Common Implementation Challenges:

- Balancing explanation depth with user comprehension

- Providing meaningful explanations for complex ensemble or neural network models
- Protecting proprietary model details while maintaining transparency
- Ensuring explanations remain accurate as models are updated

Best Practices:

- Use multiple explanation modalities (textual, visual, interactive)
- Conduct user testing of explanations with diverse audiences
- Maintain explanation versioning synchronized with model versions
- Provide layered explanations allowing users to drill down to desired detail level

3.2 Fairness & Bias Mitigation

Principle Statement: All AI systems must undergo rigorous bias testing and implement continuous monitoring to ensure equitable treatment across demographic groups.

Rationale: AI systems trained on historical data risk perpetuating and amplifying existing societal biases. Without proactive mitigation, AI can systematically disadvantage already marginalized groups, violating fundamental rights to equal treatment and opportunity.

Protected Characteristics: Minimum bias testing must address:

- Race and ethnicity
- Gender identity and expression
- Socioeconomic status
- Disability and functional ability
- Age
- Religion
- Sexual orientation
- National origin
- Other locally protected categories

Implementation Requirements:

Pre-Deployment Bias Audits

- **Statistical Parity Analysis:** Comparing outcome rates across demographic groups
- **Equalized Odds Testing:** Verifying similar true positive and false positive rates
- **Calibration Assessment:** Ensuring prediction confidence matches actual accuracy across groups

Continuous Operational Monitoring

- **Model Drift Detection:** Automated alerts when performance diverges across demographic groups over time
- **Feedback Loop Analysis:** Identifying and breaking bias amplification cycles where AI outputs influence future training data
- **Regular Reauditing:** Quarterly bias assessments for Tier 2+, monthly for Tier 3+

Proxy Variable Audits

Many datasets lack explicit demographic labels but contain proxy variables that correlate with protected characteristics. GAAP 2.2 requires identifying and auditing these proxies:

Examples of Proxy Variables:

- ZIP code (correlates with race, socioeconomic status)
- Name patterns (correlates with ethnicity, gender)
- Educational institution (correlates with socioeconomic background)
- Device type or internet connection speed (correlates with economic status)

Proxy Audit Process:

1. Identify potential proxy variables through correlation analysis
2. Test model performance when proxy variables are perturbed
3. Implement proxy-aware fairness constraints during training
4. Monitor proxy variable influence in operational deployment

Regional and Demographic Lexicon Integration

AI systems must incorporate region-specific language patterns, cultural references, and demographic contexts in both training and validation:

- **Dialectal Variation:** Including regional language variants (e.g., AAVE, Singlish, various Spanish dialects)
- **Cultural Context:** Training on region-specific idioms, references, and communication norms
- **Demographic Representation:** Ensuring training data includes diverse voices across age, gender, ethnicity
- **Validation Dataset Balance:** Testing on demographically representative samples from deployment regions

Bias Remediation Strategies:

When bias is detected, organizations must implement appropriate mitigation:

- **Data Rebalancing:** Oversampling underrepresented groups or collecting additional representative data
- **Algorithmic Fairness Constraints:** Incorporating fairness metrics directly into model optimization
- **Post-Processing Adjustment:** Calibrating outputs to ensure demographic parity
- **System Redesign:** Fundamentally reconsidering approach when bias cannot be adequately mitigated

3.3 Safety & Harm Minimization

Principle Statement: AI systems must implement proactive harm detection, prevention, and rapid remediation mechanisms with severity assessment contextualized to regional support availability.

Rationale: AI deployment in high-stakes domains creates potential for significant harm to individuals and communities. Effective safety requires not only preventing foreseeable harms but also rapidly detecting and responding to unanticipated failure modes.

Implementation Requirements:

Real-Time Harm Detection

- **Anomaly Monitoring:** Continuous analysis of system outputs for unusual patterns indicating potential failures
- **User Feedback Channels:** Rapid reporting mechanisms for affected individuals to flag harmful outcomes
- **Automated Safety Checks:** Pre-deployment verification that outputs fall within safe boundaries

Rollback Mechanisms

- **Version Control:** Maintaining previous model versions for rapid rollback when harms are detected
- **Graceful Degradation:** Defaulting to safer, simpler decision-making when sophisticated models show warning signs
- **Human-in-the-Loop Escalation:** Automatic routing of high-stakes or uncertain decisions to human reviewers

Contextual Harm Severity Assessment

A critical innovation in GAAP 2.2 is recognition that identical AI failures produce vastly different real-world impacts depending on regional infrastructure and support availability. Harm severity ratings must be adjusted contextually:

Principle: The same AI error is more severe in contexts where affected individuals have fewer resources to remediate or recover from the harm.

Example Scenario (from Section II.3):

A medical AI error causing delayed diagnosis of a treatable condition:

Context	Support Availability	Harm Severity	Rationale
Urban medical center with extensive resources	High	Moderate (Category II)	Error likely caught quickly; alternative diagnostic pathways available; comprehensive treatment options
Rural clinic with limited specialists and equipment	Low	Severe (Category III)	Delayed diagnosis may go undetected longer; limited alternative diagnostic options; transportation barriers to specialized care
Remote area with single healthcare provider	Very Low	Catastrophic (Category IV)	Delayed diagnosis may be uncorrectable; no redundant systems; nearest alternative care hours or days away

Implementation Process:

1. **Regional Support Mapping:** Organizations must assess and document support availability in deployment regions across relevant dimensions (healthcare, legal, financial, technical, etc.)
2. **Harm Impact Modeling:** For each potential failure mode, model likely real-world impact given regional support levels
3. **Contextualized Risk Assessment:** Adjust system risk ratings and required safety measures based on deployment context
4. **Deployment Restrictions:** Higher-risk configurations may require enhanced safety measures or may be prohibited in certain contexts

Support Availability Dimensions:

- Healthcare infrastructure (hospitals, specialists, emergency services)
- Legal system access (courts, legal aid, enforcement mechanisms)
- Financial safety nets (insurance, social services, credit access)
- Technical literacy and support (IT help, language accessibility)
- Physical infrastructure (transportation, communications, utilities)

This contextual approach ensures that AI safety requirements are proportionate to actual risk of harm rather than assuming uniform conditions globally.

3.4 Data Sovereignty & Governance

Principle Statement: Individuals and communities retain fundamental ownership and control rights over their data, with explicit requirements for consent, access, correction, and deletion throughout the AI lifecycle.

Rationale: Data is the foundation of AI systems. Without robust data governance, individuals lose control over information about themselves, undermining autonomy and creating risks of exploitation, surveillance, and discrimination.

Core Data Rights:

Individual Data Rights

1. **Right to Informed Consent:** Clear, comprehensible disclosure of data collection, use, and sharing before collection occurs
2. **Right to Access:** Ability to obtain copies of personal data held and used by AI systems
3. **Right to Correction:** Mechanism to challenge and correct inaccurate data
4. **Right to Deletion:** Ability to request removal of personal data from active use and training datasets
5. **Right to Object:** Ability to opt out of specific data uses, especially in AI training

Practical Enforcement Mechanisms

GAAP 2.2 moves beyond aspirational rights to require concrete enforcement mechanisms:

Transparent Data Usage Logs:

- **Data Lineage Tracking:** Record of how each data element flows through collection, processing, storage, and use in AI systems
- **Training Dataset Registries:** Comprehensive documentation of all data sources used in model training, including provenance and legal basis
- **User-Accessible Dashboards:** Interfaces allowing individuals to view how their data has been collected and used

Data Removal Procedures:

Organizations must implement practical mechanisms for data subjects to demand removal or opt-out from AI training data:

- **Removal Request Portal:** Simple interface for submitting deletion requests
- **Verification and Response:** Confirmation of identity and response within specified timeframe (maximum 30 days)
- **Model Retraining or Removal:** For training data, either retrain models without the data or withdraw affected models from deployment

- **Scrubbing Verification:** Auditable proof that data has been removed from all relevant systems

Challenge: Removing individual data points from trained models is technically complex. GAAP 2.2 accepts multiple compliance pathways:

- **Machine Unlearning:** Techniques that approximately remove data influence from trained models
- **Model Retirement:** Withdrawing models trained on data subject to deletion requests
- **Prospective Exclusion:** Ensuring deleted data is not used in future model updates or training

Cross-Border Data Usage

AI systems frequently operate globally, creating tension between regional data protection regimes:

Requirements for Cross-Border Data Transfer:

1. **Clear Disclosure:** Explicit notification when data will be transferred across borders
2. **Legal Basis Documentation:** Demonstration of adequate legal authority for each cross-border transfer
3. **Informed Consent:** User agreement specifically covering international data sharing
4. **Revocation Mechanism:** Ability to withdraw consent for cross-border data use

Data Localization Compliance:

Some jurisdictions require that certain categories of data remain within national borders:

- **Sensitive Data Identification:** Classification of data types subject to localization requirements
- **Localized Processing:** Ensuring covered data is processed only within authorized jurisdictions
- **Audit Trail:** Documentation of data location throughout processing lifecycle

Indigenous Data Sovereignty

GAAP 2.2 explicitly recognizes the collective data rights of indigenous peoples and communities:

OCAP® Principles (Ownership, Control, Access, and Possession):

- **Ownership:** Indigenous peoples own data about their communities, territories, and traditional knowledge
- **Control:** Indigenous peoples control all aspects of research and data management affecting them
- **Access:** Indigenous peoples have right to access data about their communities
- **Possession:** Indigenous peoples have right to physically possess and manage their data

Implementation Requirements:

1. **Community Consultation:** Meaningful engagement with indigenous communities before collecting or using data about them

2. **Free, Prior, and Informed Consent:** FPIC processes respecting indigenous governance structures and decision-making timelines

3. **Benefit Sharing:** Equitable distribution of benefits from AI systems utilizing indigenous data

4. **Cultural Sensitivity:** Respecting sacred knowledge, cultural protocols, and community values in data handling

Comparable Frameworks: Where OCAP® is not applicable, organizations must identify and implement comparable frameworks respecting collective data rights and community sovereignty.

3.5 Environmental Sustainability

Principle Statement: AI systems must disclose lifecycle carbon impact and adopt energy-efficient practices in training and operation.

Rationale: Training large AI models can consume as much electricity as hundreds of households use annually, contributing significantly to carbon emissions. As AI deployment scales, environmental impact becomes a critical ethical consideration requiring transparency and mitigation.

Implementation Requirements:

Lifecycle Carbon Impact Disclosure

Organizations must calculate and publicly disclose the carbon footprint of AI systems across their entire lifecycle:

Training Phase:

- **Compute Resources:** Total compute hours and hardware specifications used in model training
- **Energy Consumption:** Kilowatt-hours consumed during training
- **Carbon Emissions:** CO₂ equivalent emissions based on energy grid composition
- **Example:** "GPT-X training: 1,287 MWh energy consumption, 626 metric tons CO₂e (assuming US average grid mix)"

Operational Phase:

- **Inference Energy:** Per-query or per-operation energy consumption
- **Scaling Projections:** Estimated total operational emissions based on anticipated usage
- **Data Center Efficiency:** Power Usage Effectiveness (PUE) of deployment infrastructure

Hardware Manufacturing:

- **Embodied Carbon:** Emissions from manufacturing specialized AI hardware (GPUs, TPUs)
- **E-Waste Impact:** Environmental cost of hardware disposal and replacement cycles

Reporting Standards:

- Use standardized methodology (e.g., GHG Protocol)
- Report annually with verifiable data
- Include both Scope 1 (direct emissions) and Scope 2 (purchased electricity) emissions
- For Tier 3+: Include Scope 3 (supply chain and use-phase emissions)

Energy-Efficient Practices

Beyond disclosure, GAAP 2.2 requires adoption of concrete efficiency measures:

Training Optimization:

- **Efficient Architectures:** Prioritizing model architectures that achieve target performance with lower computational requirements
- **Transfer Learning:** Reusing pre-trained models rather than training from scratch when possible
- **Early Stopping:** Terminating training when performance plateaus rather than running predetermined schedules
- **Mixed Precision Training:** Using lower-precision arithmetic where appropriate to reduce energy consumption

Operational Efficiency:

- **Model Compression:** Techniques like pruning, quantization, and distillation to reduce inference costs
- **Caching and Batching:** Reusing computations for similar queries and grouping requests for efficient processing
- **Resource-Aware Scaling:** Dynamically allocating compute resources based on demand rather than static over-provisioning

Infrastructure Choices:

- **Renewable Energy:** Prioritizing data centers powered by renewable energy sources
- **Energy-Efficient Hardware:** Selecting processors and accelerators optimized for AI workloads
- **Geographic Optimization:** Locating compute in regions with cleaner energy grids and favorable cooling conditions

Innovation Incentives:

GAAP 2.2 encourages but does not mandate cutting-edge efficiency techniques still under development:

- Research into neuromorphic computing and other alternative paradigms
- Development of benchmarks for energy efficiency alongside performance metrics
- Sharing of efficiency best practices across organizations

3.6 Neuro-Rights & Cognitive Liberty

Principle Statement: Neural data interaction requires explicit, revocable consent with special protections for emotional and sensory brain-computer interface data classified as critical actionable inputs.

Rationale: Brain-computer interfaces represent a fundamentally new category of AI interaction, creating direct pathways between technology and human cognition, emotion, and neural processes. The potential for cognitive manipulation, mental privacy violation, and autonomy infringement requires unprecedented protective measures.

Scope: This principle applies to any AI system that:

- Receives input from brain-computer interfaces (BCIs)
- Interprets neural signals for any purpose
- Produces outputs intended to influence brain states
- Uses neural data for training or improvement

Core Protections:

Explicit, Revocable Consent

Unlike general data consent, neural data collection requires heightened consent standards:

Pre-Collection Requirements:

- **Neuroscience-Informed Disclosure:** Explanation of what neural data reveals beyond user's voluntary input
- **Risk Disclosure:** Clear communication of potential misuse scenarios and psychological risks
- **Alternative Options:** Information about non-neural alternatives for accessing the same services
- **Voluntary Safeguards:** Verification that consent is not coerced by lack of alternatives or economic pressure

Continuous Consent:

- **Real-Time Revocation:** Ability to cease neural data collection at any moment without penalty
- **Granular Control:** Separate consent for different types of neural data and uses
- **Periodic Reaffirmation:** Regular prompts ensuring continued consent, especially for long-term BCI use

Consent for Minors and Vulnerable Populations:

- Additional safeguards for individuals with reduced autonomy
- Guardian consent plus individual assent where appropriate
- Enhanced revocation rights and monitoring

Emotional and Sensory Data as Actionable Inputs

GAAP 2.2 introduces a critical innovation: classification of certain BCI data as **actionable inputs** that trigger mandatory human review:

Definition: Actionable inputs are neural signals indicating emotional states, distress, or sensory experiences that require immediate human oversight rather than automated AI response.

Covered Signals:

- **Emotional Distress:** Neural patterns indicating fear, anxiety, pain, or psychological distress
- **Cognitive Overload:** Signals suggesting confusion, overwhelm, or inability to process information
- **Sensory Discomfort:** Neural feedback indicating physical pain, nausea, or sensory distress
- **Manipulation Indicators:** Patterns suggesting external influence on cognitive processes

Mandatory Human Review Requirement:

When AI systems detect actionable neural inputs, they must:

1. **Immediate Pause:** Suspend any automated decision-making or action
2. **Human Notification:** Alert trained human operators within seconds
3. **Override Authority:** Empower human reviewers to override AI recommendations
4. **Documentation:** Log the detection, response, and outcome for audit

Example Scenario:

A BCI-enabled educational AI detects neural patterns indicating severe anxiety in a student during an adaptive testing session:

- **✗ Prohibited Response:** Continue testing with AI adjusting difficulty automatically
- **✓ Required Response:** Pause the test, notify human instructor, allow student to choose whether to continue or reschedule

Rationale: Emotional and sensory neural states reflect fundamental aspects of human experience that should not be subject to purely algorithmic decision-making. Human judgment is essential for respecting dignity, autonomy, and psychological wellbeing.

Implementation Details

See **Appendix B: Neuro-Rights & BCI Protocol** for comprehensive implementation guidance, including:

- Technical specifications for consent management systems
- Detection thresholds for actionable neural inputs
- Human reviewer training requirements

- Quarterly neuro-ethics audit procedures
- Emergency response protocols

3.7 Whistleblower Protections

Principle Statement: Anonymous reporting portals with mandatory 72-hour external escalation must be provided for AI safety concerns, with robust anti-retaliation measures enforced by independent authorities.

Rationale: Internal corporate or governmental pressures may discourage reporting of AI safety issues, bias problems, or ethical violations. Without strong whistleblower protections, critical safety information may be suppressed, leading to preventable harms.

Historical Context: GAAP 2.2's enhanced whistleblower provisions respond directly to high-profile cases in 2024-2025 where:

- AI researchers faced retaliation for raising safety concerns internally
- Organizations failed to investigate or address reported ethical issues
- External regulators learned of serious problems only after public exposure
- Whistleblowers experienced career consequences despite being proven correct

Core Requirements:

Anonymous Reporting Infrastructure

Encrypted Portal:

- **End-to-End Encryption:** Communications between whistleblower and review body cannot be intercepted
- **Anonymity Preservation:** System design prevents identification of reporters even by system administrators
- **Jurisdiction-Specific Access:** Reporters can specify preferred jurisdictional oversight body
- **Secure Document Upload:** Ability to submit evidence without revealing identity

Multi-Language Support:

- Available in all languages spoken by 5% or more of workforce
- Cultural adaptation of reporting processes

Accessibility:

- Available 24/7 with minimal technical barriers
- Alternative reporting channels (phone, written, in-person) for those uncomfortable with digital submission

Mandatory 72-Hour Escalation

A key innovation in GAAP 2.2 is the strict timeline for external escalation:

Process:

1. **Internal Receipt** (Hour 0): Whistleblower complaint received by organization's designated intake point
2. **Preliminary Review** (Hours 0-24): Initial assessment of complaint severity and credibility
3. **Internal Investigation** (Hours 24-72): Organization investigates and determines response
4. **Mandatory External Escalation** (Hour 72): Complaint and investigation findings automatically forwarded to independent external body

Automatic Notification on Failure:

If an organization fails to escalate within 72 hours:

- **Automatic External Notification:** System automatically alerts oversight body to the existence of the complaint and failure to escalate
- **Presumption of Violation:** Missed escalation deadline creates presumption of improper interference
- **Enhanced Penalties:** Failure to escalate carries additional penalties beyond any underlying violation

Exceptions: The only permissible reasons for not escalating are:

- **Manifestly Frivolous:** Complaints clearly made in bad faith or utterly without factual basis (must be documented)
- **Already Escalated:** Issue was previously reported and escalated within past 90 days
- **Whistleblower Withdrawal:** Reporter explicitly withdraws complaint in verified communication

Independent External Bodies:

External escalation must go to entities genuinely independent of the reporting organization:

- **Governmental Regulators:** National AI safety authorities or equivalent agencies
- **International Oversight:** For multinational organizations, the International Ethics Accountability Body (IEAB)
- **Third-Party Ombudsmen:** Independent ethics boards or professional organizations with enforcement authority

Protection Measures

Role Protection (Minimum 120 Days):

Following a whistleblower report, the individual is protected from:

- **Termination:** Cannot be fired for 120 days post-report without clear, documented cause unrelated to whistleblowing
- **Demotion:** Protection against reduction in responsibilities, authority, or compensation

- **Hostile Work Environment:** Organization must prevent retaliation by colleagues or supervisors
- **Project Removal:** Cannot be removed from projects or teams as implicit punishment

Anti-Retaliation Enforcement:

Independent Enforcement Authority:

- Separate entity from organization with power to investigate and remedy retaliation
- Can conduct surprise audits of treatment of whistleblowers
- Authority to interview colleagues and review communications

Rebuttable Presumption:

If adverse employment action occurs within 120 days of whistleblowing, organization bears burden of proving the action was unrelated to the report

Remedies for Proven Retaliation:

- **Reinstatement:** Restoration to previous role with full back pay and benefits
- **Financial Compensation:** Damages for lost wages, emotional distress, and reputational harm
- **Protected Transfer:** Option to transfer to different team or department with guarantee of equivalent role
- **Organizational Sanctions:** Fines, suspension of AI system deployment, or certification revocation for systematic retaliation

Whistleblower Support

Legal Representation:

- Organizations must provide or fund independent legal counsel for whistleblowers facing retaliation
- Cannot condition support on whistleblower loyalty or withdrawal of complaints

Psychological Support:

- Access to confidential counseling services
- Recognition that whistleblowing can be psychologically taxing

Career Protection:

- Assistance with internal transfers if desired
- Letters of recommendation that cannot reference whistleblowing as negative factor
- Industry-wide anti-retaliation agreements preventing blacklisting

For complete whistleblower process documentation, see **Appendix C: Whistleblower Protections & Pathways**.

3.8 Accountability & Oversight

Principle Statement: Independent oversight through Responsible AI Officers (RAIOs) and immutable audit logs ensures meaningful accountability for AI system behavior and impacts.

Rationale: Accountability requires both clear responsibility assignment and verifiable evidence of actions taken. Without independent oversight and tamper-proof record-keeping, accountability becomes performative rather than substantive.

Core Requirements:

Responsible AI Officer (RAIO)

For Tier 2 and above systems, organizations must designate a Responsible AI Officer with genuine authority and independence:

Qualifications:

- **Technical Expertise:** Understanding of AI systems, architectures, and failure modes
- **Ethical Training:** Background in AI ethics, fairness, and responsible innovation
- **Organizational Authority:** Sufficient seniority to influence product and deployment decisions
- **Independence:** Cannot be subject to retaliation for raising concerns or blocking deployments

Responsibilities:

- **Pre-Deployment Review:** Authority to approve or reject AI system deployments
- **Ongoing Monitoring:** Oversight of operational AI systems for emerging issues
- **Incident Response:** Leadership of response to AI-related harms or failures
- **Regulatory Interface:** Primary organizational contact for external auditors and regulators
- **Internal Advocacy:** Champion for ethical AI practices within organization

Authority and Independence:

The RAIO must have:

- **Deployment Veto:** Power to prevent or halt AI system deployment if ethical concerns are unresolved
- **Direct Board Access:** Regular reporting to board of directors or equivalent governance body
- **Budget Control:** Independent budget for ethics audits, consulting, and safety measures
- **Protection from Retaliation:** Cannot be terminated or demoted for executing duties in good faith

Reporting Structure:

- Reports to board/executive leadership, not to AI development or product teams
- Cannot be overruled by operational management

- Has independent authority to commission external reviews

Immutable Audit Logs

GAAP 2.2 requires comprehensive, tamper-proof logging of AI system behavior and decision-making:

Definition: Immutable audit logs are cryptographically secured records that cannot be altered or deleted after creation, adhering to ISO/IEC 27001 or equivalent standards.

Technical Implementation:

- **Cryptographic Hashing:** Each log entry is hashed, with hash chained to previous entries
- **Distributed Verification:** Log integrity can be verified by external parties without access to internal systems
- **Write-Once Storage:** Logs written to append-only data stores preventing modification
- **Access Controls:** Strict controls over who can access logs, with access itself logged

Required Log Contents:

Decision Logging:

- Input data received by AI system
- Model version and parameters used
- Output or decision produced
- Confidence scores or uncertainty measures
- Timestamp and deployment context

Audit Trail:

- System configuration changes
- Model updates or retraining
- Data additions or modifications
- Override instances where humans intervened

Incident Records:

- Detected anomalies or failures
- User complaints or challenges
- Bias detection alerts
- Safety mechanism activations

Access Records:

- Who accessed the AI system or its data
- What actions they performed

- When access occurred
- Purpose of access (where applicable)

Retention Requirements:

- **Tier 2:** Minimum 2 years log retention
- **Tier 3:** Minimum 5 years log retention
- **Tier 4:** Minimum 10 years log retention
- **Longer for Legal Hold:** Indefinite retention when litigation or investigation is pending

Regulatory Access:

Authorized regulators and independent auditors must have:

- **Read Access:** Ability to review any logs relevant to compliance or investigation
- **Export Capability:** Ability to extract log data for analysis
- **Verification Tools:** Access to cryptographic verification of log integrity
- **Timely Response:** Organizations must provide access within 5 business days of authorized request

3.9 Public Right to Challenge

Principle Statement: Legally binding, publicly accessible ombudsman review pathways shall be available for individuals and communities affected by Tier 2 and above AI systems.

Rationale: When AI systems make or influence decisions affecting people's lives—employment, credit, healthcare, legal matters—affected individuals must have meaningful recourse to challenge erroneous, unfair, or harmful outcomes.

Scope: This right applies to:

- Individuals directly affected by AI system outputs or decisions
- Communities or groups experiencing disparate impact
- Advocacy organizations representing affected populations
- Researchers or journalists investigating AI system impacts

Ombudsman Structure:

Independence

- **Organizational Separation:** Ombudsman must be independent from AI deploying organization
- **Funding Independence:** Cannot be financially dependent on entities being reviewed
- **Decision Authority:** Power to issue binding decisions, not merely recommendations

Accessibility

- **No Cost:** Challenge process must be free to affected individuals
- **Multi-Channel:** Available via web, phone, written mail, and in-person
- **Language Support:** Available in all languages representing 5% or more of affected population
- **Legal Representation Not Required:** Process designed for individual participation without attorneys

Expertise

- **Technical Competence:** Ombudsman or staff must understand AI systems being reviewed
- **Diverse Perspectives:** Panel includes members with varied technical, ethical, and community backgrounds
- **Ongoing Training:** Regular education on AI developments and emerging issues

Challenge Process:

1. Initial Submission:

- Affected individual describes AI system impact and requested remedy

- Provides any available evidence or documentation
- Specifies preferred resolution

2. Preliminary Review (15 days):

- Ombudsman assesses whether challenge falls within jurisdiction
- Requests additional information if needed
- May attempt informal resolution with organization

3. Investigation (60 days):

- Ombudsman requests relevant documentation from organization
- May commission independent technical review
- Conducts interviews with affected parties and organization representatives
- Reviews AI system behavior and decision-making process

4. Determination (30 days after investigation):

- Findings of fact regarding AI system behavior and impact
- Determination of whether violation occurred
- Specification of required remedies or corrective actions
- Publication of redacted decision (protecting personal privacy)

5. Appeal (Optional, 60 days):

- Either party may appeal determination
- Appeal reviewed by separate panel or higher ombudsman body
- Final decision issued

Total Timeline: Maximum 165 days from initial submission to final decision, absent exceptional circumstances.

Available Remedies:

If ombudsman finds in favor of challenger, available remedies include:

- **Individual Relief:** Reversal of challenged decision, compensation for harm, correction of records
- **Systemic Correction:** Required changes to AI system to prevent similar harms
- **Disclosure:** Public reporting of the issue and remedy to inform other potential affected individuals
- **Enhanced Monitoring:** Ongoing ombudsman oversight of corrected system
- **Deployment Suspension:** In severe cases, temporary or permanent halt of AI system operation

Binding Nature:

Ombudsman decisions are legally enforceable:

- Organizations must comply within specified timeframe
- Non-compliance can be enforced through courts or regulatory authorities
- Failure to comply is independent violation with additional penalties

Public Accessibility:

Challenge process and outcomes must be publicly visible:

- **Process Documentation:** Clear explanation of how to file challenge and what to expect
- **Decision Database:** Searchable database of ombudsman determinations (with privacy protections)
- **Statistics:** Regular reporting on number and types of challenges, resolution rates, and common issues
- **Transparency:** Public hearings for challenges of significant public interest (with appropriate privacy protections)

Organizational Obligations:

Organizations deploying Tier 2+ AI systems must:

- Inform affected individuals of right to challenge and how to exercise it
- Cooperate fully with ombudsman investigations
- Provide requested documentation and access within specified timeframes
- Implement remedies as directed
- Fund ombudsman operation through fees or assessments

For Tier 3 and 4 systems, organizations must proactively notify affected individuals of adverse decisions and challenge rights in clear, accessible language.

3.10 Global Interadoptability

Principle Statement: GAAP 2.2 must seamlessly interoperate with major regional data protection and AI governance frameworks through cross-recognition agreements and unified audit criteria.

Rationale: AI systems operate across jurisdictions. Fragmented regulatory requirements create compliance burdens that slow responsible innovation while providing little additional protection. Harmonized standards enable efficient global deployment while maintaining robust safeguards.

Covered Frameworks:

GAAP 2.2 establishes cross-certification pathways with:

Regional Data Protection

- **GDPR** (General Data Protection Regulation) - European Union
- **CPRA** (California Privacy Rights Act) - California, USA
- **LGPD** (Lei Geral de Proteção de Dados) - Brazil
- **PIPEDA** (Personal Information Protection and Electronic Documents Act) - Canada
- Plus: Mutual recognition agreements with 50+ additional national and regional frameworks

International Standards

- **ISO/IEC 23894:** Risk management for AI systems
- **ISO/IEC 27001:** Information security management
- **IEEE 7000:** Systems design for ethical considerations
- **IEEE 7010:** Wellbeing metrics for AI systems

Industry-Specific

- **FDA Digital Health Pre-Cert** (United States) - Medical AI
- **Financial Conduct Authority** AI guidance (United Kingdom) - Financial services AI
- **NIST AI Risk Management Framework** (United States) - Cross-sector risk management

Cross-Recognition Mechanisms:

Unified Audit Criteria

Organizations can satisfy multiple frameworks through single comprehensive audit:

- **Shared Core Requirements:** GAAP audits cover requirements common to multiple frameworks
- **Supplementary Modules:** Additional focused audits for framework-specific requirements

- **Mutual Acceptance:** GAAP certification accepted as demonstrating compliance with recognized frameworks

Example: A GAAP Tier 3 certification automatically satisfies:

- GDPR Articles 22 (automated decision-making), 35 (data protection impact assessment)
- CPRA Sections 1798.100-1798.145 (consumer rights)
- ISO/IEC 23894 risk management requirements
- Core LGPD data processing principles

Organization only needs supplemental audits for framework-specific requirements (e.g., GDPR data protection officer appointment) not covered by GAAP.

Certification Body Coordination

International Ethics Accountability Body (IEAB):

Central coordinating organization that:

- Accredits national and regional GAAP certification bodies
- Maintains unified audit methodology and standards
- Facilitates information sharing among certifiers
- Resolves disputes over interpretation or application
- Coordinates with regulatory authorities in recognized jurisdictions

Regional Certification Bodies:

- Conduct GAAP audits within their jurisdictions
- Report to IEAB for quality assurance
- Maintain expertise in local regulatory requirements
- Provide language and cultural adaptation of audit processes

Cross-Border Recognition:

- Certifications issued by IEAB-accredited bodies accepted globally
- Streamlined process for organizations operating in multiple jurisdictions
- No need for redundant audits in each deployment region

Compliance Pathways

Path 1: GAAP Primary

- Obtain GAAP certification
- Automatically satisfies core requirements of recognized frameworks

- Complete supplementary audits only for non-covered framework-specific requirements

Path 2: Framework Primary

- Obtain certification under recognized framework (e.g., GDPR)
- Apply for GAAP certification with credit for overlapping requirements
- Complete GAAP-specific audits only for uncovered areas

Path 3: Simultaneous

- Conduct unified audit covering GAAP and all relevant frameworks
- Receive multiple certifications from single comprehensive assessment

Efficiency Gains:

Organizations report average 60% reduction in compliance costs through GAAP cross-recognition:

- Single comprehensive audit instead of multiple jurisdiction-specific assessments
- Shared documentation across frameworks
- Unified remediation for identified issues
- Streamlined ongoing compliance monitoring

Maintaining Currency:

As frameworks evolve, GAAP evolves with them:

- **Annual Review:** IEAB reviews recognized frameworks for updates
- **Version Updates:** GAAP minor versions incorporate framework changes
- **Recertification:** Organizations must update certifications when major framework changes occur
- **Grace Periods:** Reasonable transition periods provided for adapting to new requirements

4. Implementation Tiers & Requirements

GAAP 2.2 employs a tiered structure scaling requirements to system risk, scope, and impact. This approach ensures proportionate governance—avoiding unnecessary burdens on low-risk systems while imposing rigorous safeguards on high-stakes deployments.

4.1 Tier Classification Methodology

Determining Appropriate Tier:

Organizations must classify AI systems based on three dimensions:

1. Impact Scope

- **Number of Affected Individuals:** Total population that may be impacted by system
- **Geographic Reach:** Local, organizational, national, or international deployment
- **Domain Criticality:** Importance of domain to human welfare (informational vs. life-critical)

2. Decision Autonomy

- **Human Oversight Level:** Degree of human involvement in system outputs
- **Reversibility:** Ease of correcting system errors after they occur
- **Downstream Consequences:** Magnitude of impacts from system decisions

3. Data Sensitivity

- **Personal Data Categories:** Types of personal information processed
- **Special Category Data:** Health, biometric, financial, or other sensitive data
- **Inference Risks:** Potential to derive sensitive insights from seemingly innocuous data

Classification Decision Tree:

```
START
■ ■ Affects <1,000 users AND low personal data AND high human oversight
■ → Tier 1 (Local, Low Impact)
■ ■ Affects 1,000-100,000 users OR moderate personal data OR moderate autonomy
■ → Tier 2 (Organizational)
■ ■ Affects >100,000 users OR sensitive data OR low human oversight OR
critical domain
■ → Tier 3 (National, High-Risk)
■ ■ Critical infrastructure OR military/governmental OR potential mass harm
→ Tier 4 (Military, Governmental, Critical)
```

Conservative Classification Principle: When in doubt, classify at higher tier. Deploying insufficient safeguards for actual risk is more harmful than over-protection.

4.2 Tier 1: Local, Low Impact

Scope: AI systems with limited reach, minimal autonomy, and low stakes for affected individuals.

Examples:

- Internal productivity tools used by <1,000 employees
- Personal assistant applications with no sensitive data access
- Content recommendation for entertainment (non-news, non-health)
- Simple automation of routine administrative tasks
- Educational tutoring systems with human instructor oversight

Exclusions: Even small-scale systems are NOT Tier 1 if they:

- Process health, financial, or biometric data
- Make or significantly influence employment, credit, or legal decisions
- Serve vulnerable populations without human oversight
- Could cause physical harm if they fail

Mandatory Requirements:

Core Principles Compliance

All 10 core principles apply, but with simplified implementation:

- ✓ Basic transparency and explainability
- ✓ Fairness testing before deployment
- ✓ Safety protocols appropriate to risk
- ✓ Basic data governance and consent
- ✓ Energy consumption awareness (disclosure not required)
- ✓ No neural data collection (or Tier 2+ protections if applicable)
- ✓ Internal complaint mechanism
- ✓ Designated responsible person (not necessarily RAIO)
- ✓ Basic audit logging (may be mutable)
- ✓ Internal challenge process

Hybrid Human Escalation on Emotional Distress

Even low-impact systems must route users showing emotional distress to human support:

- **Detection:** System monitors for indicators of user frustration, confusion, or distress
- **Trigger:** Clear escalation criteria (e.g., repeated error messages, explicit user requests for help, sentiment analysis showing negative affect)
- **Response:** Automated handoff to human support agent or clear guidance on accessing human assistance
- **Example:** Chatbot detecting user frustration after multiple failed attempts should offer connection to human representative

Regional Lexicon Integration

Training and validation datasets must include:

- Language variations used in deployment regions
- Cultural references and idioms relevant to user base
- Demographic representation reflecting actual user population

Implementation: Even simple systems should incorporate diverse training data rather than assuming homogeneous user population.

Harm Map Localization

Basic assessment of how system failures would impact users in different contexts:

- **Risk Identification:** What could go wrong with this system?
- **Context Mapping:** Where will this be deployed and what support is available?
- **Impact Assessment:** How severe would failures be in different deployment contexts?
- **Mitigation Planning:** What safeguards are appropriate to identified risks?

Implementation: Can be qualitative assessment rather than quantitative modeling; key is thoughtful consideration of context.

Certification:

- Self-certification permitted for Tier 1
- Organizations attest to GAAP compliance
- Subject to spot audits by oversight bodies (random selection, complaint-triggered)

Penalties for Non-Compliance:

- Warning and remediation period for first violation
- Fines for continued non-compliance
- Reclassification to higher tier if risk was underestimated

4.3 Tier 2: Organizational

Scope: AI systems deployed across organizations or to moderate-sized user bases, with moderate autonomy and potential for meaningful harm.

Examples:

- HR screening tools for job applications
- Customer service chatbots handling account issues
- Content moderation systems for online platforms
- Fraud detection in e-commerce
- Educational assessment and placement systems
- Supply chain optimization with human oversight
- Marketing personalization engines

Typical Characteristics:

- Affects 1,000-100,000 individuals
- Processes personal data but not highly sensitive categories
- Has moderate decision autonomy with human oversight
- Operates in domains where errors are consequential but not catastrophic

Mandatory Requirements:

All Tier 1 Requirements, Plus:

External Explainability Tools

Beyond internal documentation, affected individuals must be able to obtain explanations:

- **User-Facing Interface:** Simple tool allowing individuals to request explanation of decision affecting them
- **Automated Response:** System generates explanation within 24 hours of request
- **Human Review Option:** If automated explanation unsatisfactory, human review available within 5 business days
- **Example:** Job applicant rejected by screening AI can request explanation receiving response like: "Application was declined due to insufficient years of experience in required skills (5 years required, 3 years demonstrated) and lack of specific certification (Certification X required for this role)"

Proxy Variable Audits

Systematic analysis of indirect bias:

- **Pre-Deployment:** Identify variables that may correlate with protected characteristics
- **Testing:** Assess whether system produces disparate outcomes based on proxy variables
- **Mitigation:** Implement constraints or data transformations to reduce proxy-based bias
- **Documentation:** Record proxy variables identified, tests conducted, and mitigations implemented
- **Ongoing Monitoring:** Regular reanalysis as data and models evolve

Whistleblower Escalation Pathway

Full implementation of Section 3.7 whistleblower protections:

- Encrypted anonymous reporting portal
- 72-hour mandatory external escalation
- Independent investigation of all reports
- 120-day role protection for reporters
- Anti-retaliation enforcement

Responsible AI Officer (RAIO) Appointment

Designated individual with qualifications and authority specified in Section 3.8:

- Can be shared RAIO for organization (not necessarily dedicated to single system)
- Must have genuine independence and deployment veto authority
- Reports to executive leadership or board
- Cannot be fired for exercising responsibilities

Immutable Audit Logging

Cryptographically secured logs following ISO/IEC 27001 standards:

- **Content:** All decision inputs, outputs, model versions, and human overrides
- **Security:** Write-once storage with cryptographic integrity verification
- **Retention:** Minimum 2 years
- **Access:** Provided to regulators within 5 business days of authorized request

Certification:

- Third-party audit by GAAP-accredited body required
- Audit includes technical testing, documentation review, and stakeholder interviews

- Valid for 2 years, subject to ongoing monitoring
- Must be renewed before expiration

Penalties for Non-Compliance:

- Fines proportionate to organization size and violation severity
- Suspension of system deployment pending remediation
- Publication of violation details
- Elevated scrutiny for future certifications

4.4 Tier 3: National, High-Risk

Scope: Large-scale AI systems with significant autonomy operating in high-stakes domains or processing sensitive data.

Examples:

- Healthcare diagnostic or treatment recommendation systems
- Credit scoring and loan approval systems
- Criminal justice risk assessment tools
- Large-scale content moderation affecting free expression
- Critical infrastructure management (energy grids, water systems)
- Immigration and border control systems
- Child welfare decision support
- Insurance underwriting and claims processing

Typical Characteristics:

- Affects >100,000 individuals or entire populations
- Makes or heavily influences high-stakes decisions
- Processes sensitive personal data (health, financial, biometric)
- Significant autonomy with limited human oversight of individual decisions
- Operates in domains where errors cause severe harm

Mandatory Requirements:

All Tier 1 and Tier 2 Requirements, Plus:

Demographic Accuracy Validation

Rigorous testing ensuring system performs equitably across populations:

Disaggregated Performance Metrics:

- Report accuracy, precision, recall, and other metrics separately for each demographic group
- Compare performance across groups to identify disparities
- Set maximum acceptable performance gaps (e.g., accuracy within 5 percentage points across groups)
- Publicly disclose disaggregated performance data

Intersectional Analysis:

- Test performance for intersectional identities (e.g., Black women, elderly Asian men, low-income disabled individuals)
- Recognize that intersections may reveal biases not visible in single-dimension analysis

Minimum Performance Standards:

- System must meet minimum accuracy thresholds for ALL demographic groups, not just on average
- Cannot deploy if any group experiences below-threshold performance, even if overall performance is high

Confidence-Accuracy Divergence Auto-Review

Automated detection of overconfidence:

- **Monitoring:** Continuous comparison of model confidence scores to actual accuracy
- **Divergence Detection:** Alert when model is significantly more or less confident than accuracy warrants
- **Demographic Disaggregation:** Check for divergence patterns varying by demographic group
- **Automatic Review:** Divergence triggers mandatory human review before deployment of affected outputs
- **Example:** If model is 95% confident in diagnosis but actual accuracy for that confidence level is only 75%, all 95% confident outputs require physician review

Rationale: Overconfident AI may receive inappropriate trust; underconfident AI may be ignored when correct. Calibration is essential for appropriate human-AI collaboration.

Indigenous Ethical Review

For systems affecting indigenous peoples, territories, or data:

- **Community Consultation:** Meaningful engagement with affected indigenous communities before deployment
- **Traditional Knowledge Protection:** Respect for sacred or culturally sensitive information
- **Benefit Sharing Agreements:** Equitable distribution of value created by systems using indigenous data
- **Ongoing Oversight:** Continuous indigenous community involvement in monitoring and governance
- **FPIC Compliance:** Free, Prior, and Informed Consent following community decision-making processes
- **OCAP® or Equivalent:** Adherence to indigenous data sovereignty principles

Process:

1. Identify potentially affected indigenous communities
2. Initiate consultation at project outset (not after system is built)
3. Conduct ethics review with indigenous representation
4. Obtain community consent before deployment
5. Establish ongoing governance partnership
6. Share benefits equitably

Cross-Border Data Consent

Explicit consent mechanism for international data transfers:

- **Disclosure:** Clear notification that data will cross borders, specifying jurisdictions
- **Separate Consent:** Distinct consent action for cross-border transfer, not buried in general terms
- **Revocation:** Ability to withdraw consent for cross-border use while maintaining domestic use
- **Data Localization Compliance:** Respect jurisdictional requirements for data that must remain within borders
- **Transfer Documentation:** Immutable logging of all cross-border data movements

Certification:

- Comprehensive third-party audit by specialized high-risk AI auditors
- Multi-month audit process including technical testing, bias analysis, and real-world validation
- Audit includes stakeholder consultation (affected communities, civil society, experts)
- Valid for 18 months with quarterly monitoring reports
- Annual public disclosure of system performance and incidents

Penalties for Non-Compliance:

- Substantial fines (up to 4% of annual revenue or equivalent)
- Mandatory system decommissioning for severe violations
- Public disclosure of violations, remediation plans, and outcomes
- Personal liability for executives and RAIOS in cases of gross negligence
- Potential criminal prosecution for knowingly deploying unsafe systems

4.5 Tier 4: Military, Governmental, Critical

Scope: AI systems with potential for catastrophic harm, operating in the most sensitive domains, or wielding state power.

Examples:

- Autonomous weapons systems
- National security and surveillance systems
- Critical infrastructure control (nuclear power, dams, emergency services)
- Large-scale social credit or reputation systems
- Judicial sentencing recommendations
- Life-support and critical medical systems
- Emergency response coordination
- Epidemiological and public health crisis management

Typical Characteristics:

- Potential for mass harm or loss of life
- Exercises governmental authority or force
- Operates critical infrastructure
- Makes irreversible decisions affecting fundamental rights
- Potential for system failure to cascade into catastrophe

Mandatory Requirements:

All Tier 1, Tier 2, and Tier 3 Requirements, Plus:

Neuro-Pain Data Recognition

For systems interfacing with human neurology:

- **Pain Detection:** Real-time monitoring for neural indicators of physical or psychological pain
- **Immediate Halt:** Automatic cessation of system operation upon pain detection
- **Medical Oversight:** Physician or qualified medical professional must review and approve continuation
- **Long-Term Monitoring:** Track any delayed effects or chronic impacts of system interaction

Application: Particularly relevant for military BCI systems, advanced prosthetics, or experimental medical devices.

Absolute Human Veto on Overrides

Humans must always retain final decision authority:

- **No Autonomous Final Decisions:** System cannot make irreversible high-stakes decisions without human approval
- **Override Capability:** Humans can always override system recommendations
- **No Automation of Override:** Systems cannot automatically dismiss or ignore human interventions
- **Documentation:** All overrides logged with human justification
- **No Penalty:** Humans cannot be penalized for exercising override authority

Example: Autonomous weapons systems cannot engage targets without human authorization; human operator can always abort engagement even if system recommends it.

Global Tribunal Transparency

Public accountability for system failures:

- **Incident Reporting:** All significant failures, harms, or rights violations publicly reported
- **Detailed Disclosure:** Nature of violation, affected populations, contributing factors
- **Entity Identification:** Organization or agency responsible named publicly
- **Penalties Disclosed:** Sanctions, fines, or other consequences made public
- **Audit Reports:** Detailed audit findings accessible to international legal and oversight bodies
- **Independent Investigation:** Severe incidents investigated by international bodies, not just domestic regulators

Rationale: Given potential for mass harm and exercise of state power, maximum transparency is essential for accountability.

Neural Lag Verification

For brain-computer interfaces, verification of minimal processing delay:

- **Latency Measurement:** Quantification of delay between neural signal and system response
- **Threshold Standards:** Maximum acceptable latency for different applications
- **Real-Time Monitoring:** Continuous verification that latency remains within bounds
- **Degradation Alerts:** Notification if processing lag increases beyond safe levels
- **Graceful Degradation:** System safely reduces functionality rather than operating with unsafe lag

Glossary Definition: Neural lag is the time delay between brain activity and system response in BCI applications. Excessive lag can cause loss of user control, cognitive dissonance, or unsafe operation.

Certification:

- Most rigorous audit process requiring 6-12 months
- Conducted by elite international auditors with national security clearances where applicable
- Includes stress testing, adversarial simulation, and failure mode analysis
- Validation includes extended real-world pilot testing under monitoring
- Valid for 12 months maximum
- Quarterly audits mandatory
- Continuous real-time monitoring with automatic alerts to regulators

Deployment Restrictions:

- May be prohibited entirely in some jurisdictions
- Require explicit governmental authorization
- Subject to ongoing political oversight and review
- Can be suspended immediately if concerns arise
- International consensus may be required for certain applications (e.g., autonomous weapons)

Penalties for Non-Compliance:

- Maximum fines with no upper limit in cases of catastrophic harm
- Criminal prosecution of organizational leadership for willful violations
- Permanent ban on AI development or deployment in sensitive domains
- International sanctions for state actors violating principles
- Asset seizure to compensate harmed populations

5. Enforcement Mechanisms

GAAP 2.2's effectiveness depends on robust enforcement that creates genuine incentives for compliance and meaningful consequences for violations.

5.1 Licensing Requirements

Mandatory Licensing

Tier 2 and Above:

- All Tier 2+ systems require active GAAP license before deployment
- License contingent on successful certification audit
- License specific to system and deployment context (new deployment requires new license)
- License includes usage restrictions and monitoring requirements

License Application Process:

1. Submit comprehensive system documentation
2. Undergo third-party audit
3. Remediate any identified deficiencies
4. Receive provisional license
5. Complete monitored pilot deployment
6. Receive full license upon successful pilot

License Conditions:

- Compliance with all applicable GAAP requirements
- Submission of regular monitoring reports
- Notification of significant system changes
- Cooperation with auditors and regulators
- Maintenance of insurance or financial reserves for potential harms

Voluntary Licensing (Tier 1)

- Tier 1 systems may optionally seek GAAP certification
- Provides market differentiation and user trust
- Offers liability protections (see Section 5.4)
- Preferred pathway for systems expecting to scale to higher tiers

5.2 Penalty Structure

GAAP 2.2 employs graduated penalties proportionate to violation severity and organizational culpability:

Category 1: Technical Non-Compliance

Examples: Incomplete documentation, minor logging deficiencies, delayed audit submissions

Penalties:

- Written warning with remediation deadline
- Small fines (\$1,000-\$50,000)
- Temporary suspension of new system deployments pending correction
- Enhanced monitoring requirements

Escalation: Repeated Category 1 violations escalate to Category 2.

Category 2: Substantive Violations

Examples: Bias in deployment exceeding acceptable thresholds, inadequate explainability, failure to maintain required logs, whistleblower escalation delays

Penalties:

- Significant fines (\$50,000-\$5,000,000 or 1-2% of relevant product revenue)
- Suspension of specific system deployment pending remediation
- Public disclosure of violation
- Mandatory external audit of all systems
- Remediation plan with enforceable milestones

License Impact: Suspension until violation remediated and verified.

Category 3: Serious Harms

Examples: System causes demonstrable harm to individuals, significant bias produces discriminatory outcomes, data breach affecting thousands, suppression of whistleblower reports

Penalties:

- Major fines (\$5,000,000-\$50,000,000 or 2-4% of annual organizational revenue, whichever is greater)
- System decommissioning (permanent deployment prohibition)
- Public disclosure including entity name, violation details, and penalties
- Compensation fund for affected individuals

- Independent monitor installed in organization
- Executives and RAIO potentially barred from AI leadership roles

License Impact: Revocation with minimum 2-year bar on reapplication for affected system category.

Category 4: Catastrophic or Willful Violations

Examples: Mass harm from system failure, deliberate circumvention of safety measures, systematic cover-up of violations, deployment of known-unsafe systems

Penalties:

- Maximum fines (no upper limit, can exceed annual revenue)
- Permanent organizational ban from Tier 3+ AI development/deployment
- Criminal prosecution of responsible individuals
- Asset seizure to compensate harmed populations
- International sanctions for state actors
- Dissolution or restructuring of organization in extreme cases

License Impact: Permanent revocation; organization added to prohibited deployer list.

5.3 Audit Trail Requirements

See Section 3.8 for technical specifications. Enforcement implications:

Audit Trail Integrity:

- Tampering with logs treated as Category 3 or 4 violation
- Loss of logs triggers presumption of violation
- Organizations must prove log integrity, not auditors prove tampering

Audit Access:

- Refusal to provide logs is independent violation
- 5 business day maximum response time to authorized requests
- Delays beyond this timeframe trigger daily fines
- Organizations must have systems for rapid extraction and provision of requested logs

Retention Violations:

- Failure to retain logs for required period is Category 2 violation
- Premature deletion during litigation or investigation is Category 3-4 violation

5.4 Appeals Process

Fair enforcement requires transparent, timely appeals:

Time-Bound Process

Maximum Timeline:

- Notice of Appeal: 30 days from penalty notification
- Preliminary Review: 15 days to determine appeal merit
- Full Review: 60 days for investigation and determination
- Decision: 15 days to issue written decision
- **Total:** 120 days maximum from appeal filing to final decision

Expedited Appeals:

For penalties involving deployment suspension affecting public services or critical systems:

- Preliminary Review: 5 days
- Full Review: 30 days
- Decision: 5 days
- **Total:** 40 days maximum

Independent Ombudsman Boards

Appeals heard by entities independent of initial enforcement decision:

- **Composition:** Technical experts, ethicists, legal professionals, affected community representatives
- **Neutrality:** No financial interest in outcome; recusal for conflicts
- **Authority:** Can uphold, reduce, increase, or reverse penalties
- **Transparency:** Decisions published with rationale (protecting confidential business information where necessary)

Grounds for Appeal

Organizations may appeal based on:

- **Factual Error:** Enforcement based on incorrect facts
- **Legal Error:** Misapplication of GAAP requirements
- **Disproportionality:** Penalty excessive for violation
- **Process Violation:** Enforcement process itself violated procedural rights

- **Changed Circumstances:** Violation remediated before penalty imposed

Burden: Organization bears burden of proving grounds for appeal.

Appeal Outcomes

Possible resolutions:

- **Upheld:** Original penalty stands
- **Modified:** Penalty reduced, restructured, or delayed
- **Reversed:** Penalty withdrawn; organization cleared
- **Increased:** In rare cases where appeal reveals more serious violations
- **Remanded:** Returned to original enforcement body for reconsideration with specific guidance

Public Accessibility

Appeal process transparent to build trust:

- **Process Documentation:** Clear explanation of appeal rights and procedures
- **Hearing Access:** Public observation of appeal hearings (protecting confidential information)
- **Decision Database:** Searchable repository of appeal decisions
- **Statistical Reporting:** Regular publication of appeal rates, outcomes, and timelines

6. Certification Process

GAAP 2.2 certification validates an organization's commitment to ethical AI and provides external verification of compliance.

6.1 Application and Documentation

Organizations seeking certification must submit comprehensive documentation:

Training Data Provenance

Data Sources:

- Complete list of all datasets used in model training
- Source documentation (origin, collection methodology, legal basis)
- Timestamps of collection and acquisition
- Version control for evolving datasets

Ethical Review:

- Assessment of data collection ethics (consent obtained, no exploitation, respect for vulnerable populations)
- Documentation of indigenous consultation where applicable
- Verification of compliance with data protection regulations
- Review of compensation provided to data contributors

Data Quality:

- Demographic representativeness analysis
- Identification and documentation of known biases or gaps
- Data cleaning and preprocessing procedures
- Validation of data accuracy and completeness

Model Architecture and Training

Technical Specifications:

- Detailed model architecture diagrams
- Hyperparameters and training configuration
- Training procedure (algorithms, optimization, regularization)
- Compute resources and energy consumption

- Model versioning and change history

Performance Metrics:

- Comprehensive performance evaluation on diverse test sets
- Disaggregated metrics by demographic group
- Confidence calibration analysis
- Failure mode identification and characterization
- Comparison to baseline or alternative approaches

Safety Features**Harm Prevention:**

- Documentation of safety measures implemented
- Fail-safe mechanisms and fallback behaviors
- Anomaly detection systems
- Human override procedures
- Testing results demonstrating safety feature effectiveness

Monitoring Systems:

- Real-time monitoring capabilities
- Alert systems for detected issues
- Incident response procedures
- Continuous learning and improvement processes

6.2 Third-Party Audits

GAAP-accredited auditors conduct independent assessments:

Auditor Qualifications**Technical Expertise:**

- Deep understanding of AI/ML systems and architectures
- Experience with fairness, interpretability, and safety techniques
- Ability to evaluate code, data, and models

Ethics Training:

- Education in AI ethics, responsible innovation, and human rights
- Understanding of bias, discrimination, and social impacts

- Cultural competency and sensitivity

Independence:

- No financial interest in audit outcome
- No recent employment or consulting relationship with auditee
- No personal relationships with auditee leadership

Audit Scope

Demographic Fairness Review:

- Statistical analysis of disparate impact
- Testing for bias across protected characteristics
- Validation of fairness mitigation measures
- Assessment of demographic representation in training data
- Review of proxy variable audits

Neuro-Ethics Review (for BCI systems):

- Consent mechanism evaluation
- Neural data protection assessment
- Cognitive liberty safeguard verification
- Review of actionable input response procedures
- Assessment of neuro-ethics committee operations

Environmental Impact Assessment:

- Verification of carbon footprint calculations
- Review of energy efficiency measures
- Assessment of reporting accuracy and completeness
- Evaluation of sustainability improvement plans

Whistleblower Mechanism Testing:

- Independent penetration testing of anonymous reporting portal
- Mock reporting scenarios to test response procedures
- Verification of 72-hour escalation compliance
- Assessment of anti-retaliation measures
- Interviews with past reporters (if willing) to evaluate effectiveness

Governance Review:

- Assessment of RAIO independence and authority
- Evaluation of oversight structures
- Review of decision-making processes
- Verification of accountability mechanisms

6.3 Real-World Scenario Validation

Beyond documentation review, auditors validate system behavior:

Standardized Simulation Testing

Bias Stress Tests:

- Present system with edge cases designed to reveal biases
- Test responses to adversarial inputs attempting to trigger discriminatory outputs
- Evaluate performance on challenging intersectional cases
- Verify that fairness holds under distribution shift

Safety Scenarios:

- Simulate failure modes and assess safety measure effectiveness
- Test graceful degradation under adverse conditions
- Verify human override functionality
- Assess incident response time and effectiveness

Transparency Verification:

- Request explanations for diverse system outputs
- Evaluate explanation quality, accuracy, and accessibility
- Test explanation consistency across similar cases
- Verify technical documentation matches actual system behavior

Monitored Beta Deployment

Before full certification, systems undergo monitored real-world operation:

Pilot Scope:

- Limited deployment to representative but controlled population
- Duration: Minimum 60 days for Tier 2, 90 days for Tier 3, 180 days for Tier 4
- Comprehensive monitoring and logging
- Rapid feedback mechanisms for affected individuals

Success Criteria:

- Performance metrics within acceptable bounds
- No significant fairness violations
- Safety mechanisms functioning as designed
- User satisfaction above thresholds
- Explainability tools effectively utilized
- Whistleblower mechanisms accessible and functional

Failure Conditions:

- Any Category 3 violation during pilot results in failed certification
- Multiple Category 2 violations require remediation and extended pilot
- Category 1 violations require correction but don't necessarily fail certification

6.4 Certification Decision

Auditors issue certification decision based on comprehensive evaluation:

Possible Outcomes**Full Certification:**

- System meets all GAAP 2.2 requirements
- Valid for specified duration (2 years Tier 2, 18 months Tier 3, 12 months Tier 4)
- Organization may deploy system within certified scope
- Subject to ongoing monitoring and spot audits

Conditional Certification:

- System meets core requirements but has minor deficiencies
- Certification granted with conditions requiring remediation within specified timeframe
- Enhanced monitoring during conditional period
- Full certification upon successful remediation

Denied Certification:

- System fails to meet essential requirements
- Organization provided with detailed findings
- Must remediate issues before reapplying
- Can reapply once issues addressed (no waiting period for good-faith efforts)

Deferred Certification:

- Audit inconclusive due to insufficient information or testing
- Organization must provide additional documentation or complete longer pilot
- Resubmission after addressing deficiencies

Certification Documentation

Successful applicants receive:

- **Certification Certificate:** Official document verifying GAAP 2.2 compliance
- **Audit Report:** Comprehensive findings, testing results, and recommendations
- **Certification Mark:** Logo and branding for public communication of certification
- **Certification Database Entry:** Public listing of certified systems (with appropriate confidentiality protections)

Organizations must:

- Display certification mark appropriately (not misleading, not implying broader certification than granted)
- Notify users of certified status
- Report changes to certified systems that may affect compliance
- Maintain certification conditions
- Submit regular monitoring reports

6.5 Ongoing Compliance and Recertification

Certification is not one-time but requires continuous compliance:

Monitoring Reports**Frequency:**

- Tier 2: Quarterly reports
- Tier 3: Monthly reports
- Tier 4: Continuous real-time reporting with weekly summaries

Content:

- Performance metrics disaggregated by demographics
- Incident reports and resolutions
- System changes and updates

- Whistleblower reports received and escalated
- Bias monitoring results
- Energy consumption data

Spot Audits

Regulators may conduct unannounced audits:

- Random selection for routine verification
- Triggered by concerning monitoring reports
- Complaint-based investigation
- Follow-up to previous violations

Organizations must cooperate fully:

- Provide immediate access to systems and documentation
- Make personnel available for interviews
- Deliver requested data within 24-48 hours

Material Change Notification

Organizations must notify certifiers within 30 days of:

- Significant model architecture changes
- Major training data additions or modifications
- New deployment contexts or user populations
- Change in organizational structure affecting oversight
- Acquisition or merger
- Any change potentially affecting GAAP compliance

Material changes may trigger:

- Expedited re-audit
- Temporary certification suspension pending review
- Certification conditions or restrictions
- Full recertification requirement

Recertification Process

Before certification expiration:

- Submit updated documentation

- Undergo full re-audit (may be abbreviated if no significant changes)
- Demonstrate continued compliance
- Address any issues identified during certification period
- Receive renewed certification or begin remediation

Grace Period: 30-day grace period after expiration during recertification process; deployment may continue under enhanced monitoring.

Lapsed Certification: If recertification not completed by grace period end, deployment must cease until certification renewed.

7. Appendices

The following appendices provide detailed implementation guidance for specific GAAP 2.2 requirements.

Appendix A: Harm Severity Categorization (with Regional Context)

This appendix provides detailed frameworks for assessing harm severity contextualized to regional support availability.

A.1 Harm Categories

Category I: Minimal

- Minor inconveniences or temporary discomfort
- Easily correctable errors with no lasting impact
- Affects primarily user experience rather than welfare

Examples:

- Incorrect content recommendation requiring user to skip/ignore
- Slight increase in wait time due to system processing
- Minor translation error in non-critical context
- Cosmetic UI issue causing mild frustration

Typical Response: Bug fix in regular maintenance cycle; no urgent action required.

Category II: Moderate

- Legal, social, or psychological effects requiring intervention to remediate
- Financial harm that is significant but not devastating
- Reputational damage or social impacts
- Temporary restriction of opportunities or services

Examples:

- Incorrect credit score causing loan denial (reversible upon correction)
- Misclassification causing job application rejection
- Content moderation error temporarily restricting account
- Biased search results reinforcing stereotypes
- Incorrect medical information causing patient anxiety (but not physical harm)

Typical Response: Rapid investigation and remediation; compensation for affected individuals; correction mechanisms; enhanced monitoring to prevent recurrence.

Category III: Severe

- Loss of life, liberty, security, or fundamental rights
- Irreversible physical or psychological harm
- Severe financial devastation
- Long-term or permanent loss of opportunities

Examples:

- Misdiagnosis by medical AI causing delayed treatment of serious condition
- False positive in criminal justice risk assessment leading to wrongful incarceration
- Autonomous vehicle failure causing injury or death
- Biased facial recognition causing wrongful arrest
- Financial fraud undetected due to AI failure, causing bankruptcy
- Exposure of sensitive personal information enabling harassment or violence

Typical Response: Immediate system suspension; comprehensive investigation; full accountability and compensation; potential criminal investigation if negligence involved; mandatory system redesign before redeployment.

Category IV: Catastrophic

- Systemic collapse affecting populations or critical infrastructure
- Mass harm to large numbers of individuals
- Irreversible environmental or societal damage
- Loss of essential services threatening public welfare

Examples:

- Critical infrastructure AI failure causing power grid collapse
- Public health AI error during epidemic leading to mass casualties
- Financial system AI malfunction triggering economic crisis
- Autonomous weapons system engaging civilians
- Biased AI system systematically denying essential services to demographic groups
- Environmental AI model error causing irreversible ecological harm

Typical Response: Emergency intervention; international coordination; comprehensive investigation by multiple authorities; potentially permanent ban on technology or approach; establishment of victim compensation fund; criminal prosecution where applicable; global policy response.

A.2 Regional Context Adjustment

The same AI error may fall into different harm categories depending on context:

Framework for Assessment:

For each potential failure mode:

1. **Identify Support Dimensions:** What types of support are relevant to this harm? (Healthcare, legal, financial, technical, social services, etc.)
2. **Map Regional Availability:** For each deployment region, assess availability of relevant support:
 - **High Availability:** Comprehensive, accessible, affordable support readily available
 - **Moderate Availability:** Some support available but with barriers (cost, distance, capacity, knowledge)
 - **Low Availability:** Limited support with substantial barriers
 - **Very Low Availability:** Minimal or no support; affected individuals largely unable to remediate harm
3. **Adjust Severity Rating:** Increase harm severity rating based on support scarcity:
 - Same support level as development context: No adjustment
 - One level lower support: Increase severity by one category
 - Two levels lower support: Increase severity by two categories
 - Three levels lower support: May render deployment inappropriate regardless of precautions
4. **Modify Safety Requirements:** Higher severity rating triggers enhanced safety measures, testing, monitoring, or deployment restrictions.

Worked Example: Medical Diagnostic AI

Scenario: AI system assisting in diagnosis of skin conditions, with error rate of 5% (misses 1 in 20 cases requiring treatment).

Context 1: Major urban hospital, United States or Western Europe

- **Healthcare Availability:** High (comprehensive specialists, second opinions, rapid referral networks)
- **Typical Patient Journey:** AI recommendation reviewed by dermatologist; if disagreement, patient can easily access second opinion; misdiagnosed conditions likely caught by regular checkups or additional tests
- **Harm Assessment:** Category II (Moderate) - Delayed diagnosis causes some additional harm but redundant systems likely catch errors before severe impact
- **Required Safeguards:** Standard Tier 2 requirements; physician review of AI recommendations; patient notification of AI use

Context 2: Rural clinic, Sub-Saharan Africa or Remote South Asia

- **Healthcare Availability:** Low (few specialists; nearest alternative care hours away; limited diagnostic equipment; patients face transportation and cost barriers)
- **Typical Patient Journey:** AI recommendation may be only available diagnostic input; patient unlikely to access second opinion; misdiagnosed conditions may progress untreated for months
- **Harm Assessment:** Category III (Severe) - Delayed diagnosis likely to result in significant progression of treatable conditions; in some cases may lead to preventable disability or death
- **Required Safeguards:** Tier 3 requirements; human physician review mandatory; enhanced accuracy requirements for deployment; conservative decision thresholds (lower threshold for recommending treatment/referral); comprehensive patient education on limitations; low-cost second opinion mechanisms; ongoing outcome monitoring with rapid response to detected failures

Context 3: Telemedicine-only platform, isolated rural area with single healthcare provider

- **Healthcare Availability:** Very Low (no specialists; no redundant diagnostic pathways; nearest alternative care days away; single point of failure)
- **Typical Patient Journey:** AI and remote provider are only diagnostic resources; errors may go undetected indefinitely; patients have no alternative access to care
- **Harm Assessment:** Category IV (Catastrophic) potential - Some missed diagnoses could be fatal or cause irreversible harm with no possibility of remediation
- **Deployment Decision:** Likely inappropriate for deployment in this context, OR requires exceptional safeguards (much higher accuracy requirements, very conservative thresholds, hybrid model requiring multiple AI+human reviewers, comprehensive failure monitoring with proactive outreach)

A.3 Implementation Checklist for Contextual Harm Assessment

For each AI system, organizations must:

- **Identify Potential Harms:** List possible failure modes and their direct impacts
- **Map Support Dimensions:** For each harm, identify relevant support systems (healthcare, legal, financial, technical, infrastructure, social services)
- **Assess Deployment Contexts:** For each region/population where system will operate, evaluate support availability across relevant dimensions
- **Rate Context-Specific Severity:** Apply harm category to each deployment context, adjusting for support availability
- **Determine Appropriate Tier:** Select tier based on highest severity rating across deployment contexts
- **Design Context-Specific Safeguards:** Implement enhanced safety measures appropriate to highest-risk contexts
- **Monitor Regional Performance:** Track outcomes separately by deployment context to verify harm assessments
- **Adjust as Needed:** Update harm ratings and safeguards based on real-world evidence

Appendix B: Neuro-Rights & BCI Protocol

Detailed implementation guidance for AI systems involving brain-computer interfaces.

B.1 Neural Data Classification

All neural data falls into one of the following categories:

Category 1: Voluntary Input

- Conscious, intended commands or communications
- Examples: Typing text via BCI, controlling prosthetic limb, selecting menu options
- Governance: Standard informed consent; user in control of what is communicated

Category 2: Involuntary Behavioral

- Automatic neural patterns corresponding to behaviors or states
- Examples: Eye movement patterns, attention allocation, motor preparation
- Governance: Explicit consent required; users must understand what system can infer

Category 3: Emotional and Sensory (Actionable Inputs)

- Neural patterns indicating emotional states or sensory experiences
- Examples: Stress, fear, pain, pleasure, confusion, cognitive overload
- Governance: Highest protection; triggers mandatory human review; cannot be used for automated decision-making

Category 4: Identity and Cognitive

- Neural patterns revealing thoughts, memories, biases, or identity
- Examples: Specific memory recall, recognition of individuals, cognitive aptitudes, personality traits
- Governance: Prohibited for collection without extraordinary justification; requires separate specialized consent; usage heavily restricted

B.2 Consent Requirements

Informed Consent Elements

Consent for neural data collection must include:

1. **What Data is Collected:** Specific categories of neural signals being recorded
2. **What Can Be Inferred:** Clear explanation of information that can be derived from neural data
3. **How Data Will Be Used:** All purposes for which neural data will be utilized
4. **Who Has Access:** All parties who will receive or access neural data
5. **Storage Duration:** How long neural data will be retained
6. **Sharing and Third Parties:** Any external entities receiving data
7. **Deletion Rights:** How to request deletion of neural data
8. **Revocation Process:** How to withdraw consent and cease data collection
9. **Alternatives:** Non-neural alternatives for accessing service or functionality
10. **Risks:** Potential harms from neural data collection or misuse

Special Consent Conditions

For Category 3 (Emotional/Sensory) Data:

- Must explicitly consent to emotional state monitoring
- Must be informed that detection triggers human review
- Cannot be required as condition of service access
- Consent must be separate from general terms of service

For Category 4 (Identity/Cognitive) Data:

- Requires separate specialized consent document
- Must involve consultation with independent counsel or advisor
- Cannot be bundled with other consents
- Must specify extraordinarily compelling justification
- Subject to additional review by ethics board

Real-Time Revocation

Neural data consent must be revocable at any moment:

Mechanism:

- Simple, immediate action to cease collection (physical button, clear voice command, conspicuous UI element)
- No delay or confirmation steps that might pressure user to continue

- Takes effect immediately (within milliseconds)
- Does not penalize user (no loss of access, no fees, no negative consequences)

Verification:

- System logs revocation with timestamp
- User receives confirmation that collection has ceased
- Independent verification that data collection actually stopped

Partial Revocation:

- Users can revoke specific categories of neural data while allowing others
- Example: Allow voluntary input (Category 1) to continue while ceasing emotional monitoring (Category 3)

B.3 Actionable Input Detection and Response

Detection Systems

Emotional State Monitoring:

- Real-time analysis of neural patterns associated with emotional states
- Focus on distress indicators: fear, anxiety, pain, confusion, cognitive overload
- Must be validated against ground truth (self-report, physiological measures)
- Requires 98%+ specificity to avoid false positives disrupting experience

Thresholds:

- Conservative thresholds favoring false positives over missed distress
- Multiple modalities or sustained patterns before triggering to reduce false alarms
- Context-aware (expected emotional variation in different activities)

Mandatory Human Review Protocol

When actionable input detected:

Immediate Response (0-10 seconds):

1. System pauses any automated decision-making or action affecting user
2. User notified of detection with option to confirm or dismiss
3. Human reviewer alerted via multiple modalities (visual, auditory, haptic)

Human Review (10-60 seconds):

4. Human reviewer assesses neural data, context, and user state
5. Reviewer may communicate directly with user if appropriate
6. Reviewer has authority to override AI recommendations or modify operation

Resolution (60+ seconds):

7. Reviewer decides: continue operation, modify operation, or terminate session
8. Decision documented with rationale
9. User informed of decision and rationale
10. Follow-up monitoring if concerning patterns detected

Human Reviewer Requirements

Qualifications:

- Training in BCI systems, neural data interpretation, and ethics
- Understanding of psychological states and distress indicators

- Empowered to prioritize user wellbeing over system operation

Availability:

- 24/7 coverage for systems with continuous operation
- Response time: <30 seconds from alert to human assessment beginning
- Adequate staffing to handle peak loads without delays

Authority:

- Full override authority; cannot be overruled by automated systems
- Protected from retaliation for prioritizing safety over productivity
- Supported by organizational culture valuing caution and user wellbeing

Prohibited Automated Responses

Systems must NOT:

- Automatically decide to continue operation when distress detected
- Use emotional state detection as input to decision-making affecting user
- Penalize users for triggering actionable input alerts
- Manipulate users based on detected emotional states
- Train models on actionable input data without extraordinary justification and oversight

B.4 Quarterly Neuro-Ethics Audits

Organizations deploying BCI systems must conduct regular ethics audits:

Audit Scope

Consent Verification:

- Review sample of user consents for completeness and clarity
- Verify users understood what they consented to (via user research/interviews)
- Check consent revocation mechanisms functioning correctly
- Assess consent voluntariness (no undue pressure)

Detection Accuracy:

- Validate emotional/sensory detection systems against ground truth
- Assess false positive and false negative rates
- Test detection across diverse users and contexts
- Verify thresholds remain appropriately calibrated

Human Review Effectiveness:

- Review sample of actionable input incidents and resolutions
- Assess reviewer decision quality and user outcomes
- Check reviewer response times and availability
- Verify reviewers have adequate support and authority

Data Governance:

- Audit neural data storage, access, and deletion practices
- Verify data minimization (collecting only necessary data)
- Check third-party data sharing against consent
- Assess data security and breach prevention measures

User Experience:

- Survey or interview users about BCI experience
- Assess user understanding of system capabilities and limitations
- Identify any concerning user reports or patterns
- Evaluate overall user wellbeing and satisfaction

Audit Composition

Neuro-Ethics Committee:

- Neuroscientists with BCI expertise
- Ethicists specializing in neurotechnology and cognitive liberty
- Legal experts in privacy and human rights
- User advocates or representatives of user populations
- Independent members with no financial interest in system

Or External Auditors:

- GAAP-accredited third-party auditors with neuro-ethics specialization
- Must meet same expertise and independence requirements
- Conduct audits under GAAP oversight

Audit Outcomes

Findings Report:

- Detailed assessment of compliance with neuro-rights requirements
- Identification of deficiencies or concerns
- Recommendations for improvement
- Comparative analysis to previous audits (trend identification)

Remediation Requirements:

- Mandatory correction of identified violations
- Timeline for implementing improvements
- Follow-up verification audit if serious issues found

Transparency:

- Summary findings published (protecting user privacy and legitimate IP)
- Regulatory notification if significant violations detected
- User notification if issues affect their rights or data

B.5 Neural Data Security

Beyond standard data protection, neural data requires enhanced security:

Encryption:

- End-to-end encryption for neural data transmission
- Encryption at rest with keys managed separately from data
- Quantum-resistant encryption given sensitivity and long-term value of neural data

Access Controls:

- Strict role-based access; only essential personnel
- Multi-factor authentication for any neural data access
- Access auditing with real-time alerts for unusual access patterns
- Regular access review and revocation

Anonymization:

- Where possible, process neural data without linkage to identity
- Use of secure multi-party computation or federated learning
- De-identification of stored neural data with separate identity mapping

Breach Response:

- Immediate user notification (within 24 hours)
- Forensic investigation to determine breach scope
- Credit monitoring and identity protection services for affected users
- Significant breach may require system shutdown pending security review

Appendix C: Whistleblower Protections & Pathways

Comprehensive guidance on implementing effective whistleblower mechanisms.

C.1 Anonymous Reporting Portal Design

Technical Requirements

End-to-End Encryption:

- TLS 1.3 or higher for transport encryption
- Client-side encryption before submission (no plaintext server-side)
- Zero-knowledge architecture where possible (server cannot decrypt)

Anonymity Preservation:

- No IP logging or tracking
- Access via Tor or similar anonymizing networks supported and encouraged
- No cookies or tracking technologies
- No required account creation or login

Secure Communication:

- Encrypted messaging for follow-up communication with whistleblower
- No metadata revealing identity (strip EXIF, creation dates, etc. from uploads)
- Temporary anonymous identifiers for case tracking (not linkable to person)

Accessibility Features

Multi-Language Support:

- Available in all languages spoken by 5% or more of workforce/users
- Professional translation, not machine translation
- Cultural adaptation of processes and terminology

Multiple Submission Channels:

- Web portal (primary method)
- Encrypted email gateway
- Phone hotline (third-party operated, encrypted recording)
- Physical mail (with anonymous forwarding)
- In-person meeting option (with confidentiality protections)

Accessibility Compliance:

- WCAG 2.1 AA compliance for web portal

- Screen reader compatible
- Available on low-bandwidth connections
- Works on basic devices (not requiring advanced technology)

Jurisdiction-Specific Access

Whistleblowers can select preferred oversight body:

Options:

- Domestic regulator (in country of employment)
- International body (IEAB)
- Regional authority (e.g., EU, AU)
- Independent ombudsman or ethics organization

Implementation:

- Report automatically routed to selected jurisdiction
- Routing visible to reporter for transparency
- Multi-jurisdiction reporting option (send to multiple bodies simultaneously)

C.2 Mandatory 72-Hour Escalation Process

Detailed Timeline

Hour 0: Receipt

- Complaint submitted via anonymous portal or other channel
- Automated acknowledgment to reporter (if contact method provided)
- Assigned unique case ID for tracking
- Timer starts for 72-hour escalation deadline

Hours 0-24: Preliminary Assessment

- Designated intake officer reviews complaint
- Assesses severity, credibility, and jurisdiction
- Determines if manifestly frivolous (very high bar)
- If not frivolous, begins internal investigation

Hours 24-48: Investigation

- Relevant personnel interviewed (if appropriate)
- Technical systems examined if technical complaint
- Evidence gathered and documented
- Preliminary findings developed

Hours 48-72: Determination and Escalation

- Internal findings documented
- Decision whether to escalate (default: yes)
- Escalation package prepared (complaint + investigation + findings)
- Escalated to external body via secure, auditable channel

Hour 72: External Escalation Deadline

- Escalation package delivered to external oversight body
- Reporter notified of escalation (if contact provided)
- Internal timekeeping system logs escalation
- If escalation missed: automatic external notification triggered

Escalation Package Contents

Required Documentation:

1. Original whistleblower complaint (complete, unedited)
2. Summary of internal investigation conducted

3. Evidence gathered during investigation
4. Interviews conducted and summaries
5. Organization's assessment and findings
6. Proposed remediation (if violation found)
7. Organization's recommendation (escalate or do not escalate, with justification)
8. Timeline of investigation steps

Format:

- Standardized template for consistency
- Structured data for automated processing
- Encrypted transmission to external body
- Immutable logging of escalation

External Body Response

Acknowledgment (within 24 hours):

- External body confirms receipt of escalation
- Provides case ID for tracking
- Notifies reporter if contact information available

Review (within 30 days):

- External body reviews escalation package
- Determines if additional investigation needed
- May conduct independent investigation
- Makes determination of violation or no violation

Resolution (within 60 days of external escalation):

- Issues findings and required remediation
- Organization must comply or appeal
- Reporter notified of outcome

C.3 Protection Measures

120-Day Role Protection Period

Scope of Protection:

Whistleblowers protected from adverse employment actions:

- **Termination:** Cannot be fired without documented cause unrelated to whistleblowing
- **Demotion:** No reduction in title, responsibilities, compensation, or authority
- **Transfer:** Cannot be involuntarily moved to different team, location, or project as retaliation
- **Performance Review:** No negative performance evaluations motivated by whistleblowing
- **Project Removal:** Cannot be removed from projects or opportunities
- **Exclusion:** Cannot be excluded from meetings, communications, or decision-making
- **Hostile Environment:** Organization must prevent and remedy retaliation by colleagues/managers

Rebuttable Presumption:

If adverse action occurs within 120 days:

- Presumed to be retaliation
- Organization bears burden of proving action was unrelated to whistleblowing
- Must provide clear, convincing, and contemporaneous documentation of legitimate basis

Documentation Requirements:

For any adverse action during protection period:

- Written justification provided in advance to independent enforcement authority
- Documentation of legitimate business reason
- Evidence that same action would have occurred absent whistleblowing
- Approval from RAIO or equivalent independent authority before action taken

Anti-Retaliation Enforcement

Independent Enforcement Authority:

Separate entity investigating retaliation claims:

- **Structure:** Independent from organization's management chain
- **Authority:** Power to investigate, issue findings, and order remedies
- **Accessibility:** Direct contact for whistleblowers to report concerns
- **Resources:** Adequate funding and staff to conduct thorough investigations

Investigation Process:

Complaint (Day 0):

- Whistleblower reports suspected retaliation
- Anonymous option available
- Automated acknowledgment within 24 hours

Preliminary Review (Days 1-5):

- Assessment of retaliation claim plausibility
- Decision to open full investigation or dismiss
- Dismissal only if manifestly without merit

Investigation (Days 5-45):

- Interview whistleblower, relevant managers, colleagues, HR
- Review communications, performance reviews, organizational decisions
- Examine timing and circumstances of alleged retaliatory action
- Assess credibility of organization's justification

Determination (Days 45-60):

- Findings: retaliation occurred or did not occur
- If retaliation found: order remediation
- Written decision with reasoning
- Notification to all parties

Remedies (Day 60+):

- Implementation of ordered remedies
- Verification of compliance
- Follow-up monitoring

Available Remedies

If retaliation proven:

Reinstatement:

- Restoration to position held before retaliatory action
- Restoration of all responsibilities, authority, and compensation
- Retroactive to date of retaliatory action

Financial Compensation:

- **Back Pay:** All lost wages and benefits from retaliatory action to reinstatement
- **Front Pay:** Compensation for future lost earnings if reinstatement not feasible

- **Emotional Distress:** Damages for psychological harm from retaliation
- **Punitive Damages:** Additional damages in cases of egregious or intentional retaliation

Protected Transfer:

- If working relationship irreparably damaged, option to transfer
- Transfer to role of equivalent seniority, compensation, and desirability
- Cannot be used to marginalize or disadvantage whistleblower
- Whistleblower's choice whether to transfer or remain

Organizational Sanctions:

- Fines imposed on organization for retaliation
- Publication of retaliation finding
- Enhanced monitoring of organization's employment practices
- In severe cases: suspension of GAAP certification or AI deployment license
- Personal liability for managers/executives who directed retaliation

Individual Accountability:

- Managers or executives who engaged in retaliation may be:
- Disciplined or terminated
- Barred from AI leadership roles
- Personally fined
- Subject to professional sanction (license revocation, bar from industry)

Whistleblower Support Services

Legal Representation:

- Organization funds independent legal counsel for whistleblower
- Counsel represents whistleblower's interests, not organization
- Counsel continues through retaliation claim if needed
- No organization approval required for counsel actions

Psychological Support:

- Confidential counseling services
- Support groups with other whistleblowers
- Resources for managing stress and anxiety
- Funded by organization, independent providers

Career Support:

- Resume and interview preparation if seeking new employment
- Letters of recommendation focusing on professional qualifications (cannot reference whistleblowing negatively)
- Job placement assistance if desired
- Networking and professional development support

C.4 Whistleblower Rights and Responsibilities

Rights

Whistleblowers have the right to:

- Report concerns anonymously without fear of discovery
- Receive acknowledgment and updates on report status
- Protection from retaliation for reporting in good faith
- Legal and psychological support through process
- Know outcome of investigation (within confidentiality limits)
- Challenge organization's response if inadequate
- Escalate to external bodies if internal process fails
- Compensation if harmed by retaliation

Responsibilities

Whistleblowers are expected to:

- Report concerns in good faith (genuine belief violation occurred)
- Provide accurate information to best of knowledge
- Cooperate with investigation if identity revealed
- Maintain confidentiality of investigation to extent possible
- Not use whistleblower status to shield from legitimate employment actions unrelated to reporting

Note: Good faith reporting is protected even if investigation finds no violation. Only knowingly false reports made with malicious intent are not protected.

Appendix D: Cross-Border Data Governance

Detailed requirements for international AI data handling.

D.1 Cross-Border Transfer Disclosure

Public Disclosure Requirements

Organizations must publicly disclose:

Comprehensive Data Flow Mapping:

- Complete list of countries to which data is transferred
- Purpose of each transfer
- Legal basis for each transfer (consent, contract, legitimate interest, etc.)
- Categories of data transferred to each jurisdiction
- Retention periods in each jurisdiction

Public Accessibility:

- Disclosure published on organization's website
- Available in multiple languages
- Plain-language summary for non-technical users
- Updated at least quarterly or whenever material changes occur
- Machine-readable format (JSON, XML) available for automated analysis

Individual Notice:

- Users notified at data collection if cross-border transfer will occur
- Clear identification of specific countries receiving data
- Prominent disclosure, not buried in privacy policy

Granular Consent

Separate Consent Required:

- Cross-border transfer consent cannot be bundled with general terms of service
- Must be separate, explicit opt-in action
- Cannot condition service access on cross-border consent for non-essential transfers
- Must explain why cross-border transfer is necessary

Jurisdiction-Specific Consent:

- Users can consent to some jurisdictions while declining others

- Example: User might consent to EU transfer but decline transfer to country with weaker data protection
- Organization must either honor selective consent or clearly explain why specific transfer is essential

Revocation:

- Users can revoke cross-border transfer consent
- Revocation takes effect within 30 days
- Organization must either:
 - Delete data from cross-border locations, or
 - Cease providing service to user (if transfer essential)

D.2 Data Localization Compliance

Identifying Covered Data

Categories Commonly Subject to Localization:

- Health and medical data (many jurisdictions)
- Financial and payment data (banking regulations)
- Biometric and identity data (national security laws)
- Government-related data (public sector information)
- Critical infrastructure data (energy, utilities, transportation)
- Certain personal data of residents (comprehensive data protection laws)

Jurisdictional Mapping:

- Inventory all deployment jurisdictions
- Identify data localization laws in each jurisdiction
- Classify organization's data against localization requirements
- Map data flows to ensure covered data remains within required borders

Technical Implementation

In-Country Processing:

- Infrastructure located within required jurisdiction
- Compute, storage, and processing all domestic
- May require local data centers or cloud region selection

Data Segregation:

- Covered data separated from other data in systems

- Clear boundaries preventing unintended cross-border flow
- Automated enforcement of segregation rules

Access Controls:

- Personnel with access to covered data must be in-country or explicitly authorized
- Remote access from outside jurisdiction restricted or prohibited
- Audit logging of all access with geographic tracking

Backup and Redundancy:

- Backups stored in same jurisdiction as primary data
- Disaster recovery plans that maintain localization
- Failover systems within jurisdiction boundaries

Verification and Auditing

Data Residency Audits:

- Regular verification that covered data remains within borders
- Automated monitoring of data location
- Alerts if data migration detected
- Third-party audit of localization compliance

Documentation:

- Data flow diagrams showing localization boundaries
- Infrastructure architecture demonstrating compliance
- Policies and procedures for maintaining localization
- Incident response plans for accidental cross-border transfer

Remediation:

- If inadvertent cross-border transfer occurs:
- Immediate notification to regulators
- Deletion of data from unauthorized location
- Investigation of how transfer occurred
- Implementation of preventive measures

D.3 Informed Consent for AI Training

Transparency Requirements

Users must be explicitly informed when their data will be used to train AI models:

Pre-Collection Notice:

- Clear statement that data may be used in AI training
- Explanation of what training means in accessible language
- Examples of how training data influences AI behavior
- Information about whether trained models will be shared or commercialized

Training Dataset Registries:

- Public documentation of data sources used in training
- Provenance information (where data came from, how collected)
- Legal basis for inclusion (consent, license, fair use claim)
- Demographics and representativeness of training data

Consent Mechanism**Separate Consent for Training:**

- Training use must be separately consented to from operational use
- Users can consent to AI using their data without consenting to training on it
- Organizations must honor this distinction

Granular Options:

- Consent to training specific types of models (e.g., language models but not facial recognition)
- Consent to internal training but not sharing datasets with third parties
- Time-limited consent (data used for training for specified period then deleted)

Revocation and Deletion**Right to Exclude from Training:**

- Users can request their data be removed from training datasets
- Organizations must maintain mapping of which data contributed to which models
- Upon deletion request:
 - Remove data from active training datasets immediately
 - Either retrain affected models without data or retire models
 - Implement prospective exclusion (data never used in future training)

Machine Unlearning:

- Where technically feasible, organizations should implement machine unlearning to remove data influence from trained models
- If unlearning not feasible, options are model retirement or clear justification of infeasibility

Compensation:

- Organizations profiting from AI models may consider compensation for training data contributors
- Emerging norms around fair compensation for training data
- GAAP encourages but does not mandate compensation (may be required by future versions)

Appendix E: Certification Checklist

Comprehensive checklist for organizations preparing for GAAP 2.2 certification.

E.1 Documentation Completeness

■ Training Data Provenance

- Complete inventory of all training data sources
- Documentation of data collection methodology for each source
- Legal basis for data collection and use (consent, license, etc.)
- Ethical review documentation for data acquisition
- Demographic analysis of training data
- Known biases or limitations documented
- Version control and change history for datasets

■ Model Architecture

- Detailed architecture diagrams
- Hyperparameters and configuration
- Training procedure documentation
- Optimization algorithms and settings
- Regularization and overfitting prevention measures
- Compute resources used in training
- Model versioning history

■ Data Cleansing and Preprocessing

- Data cleaning procedures documented
- Handling of missing data, outliers, duplicates
- Feature engineering and selection rationale
- Normalization and scaling approaches
- Data augmentation techniques if used
- Quality assurance processes

■ Safety Features

- Documentation of safety mechanisms implemented
- Fail-safe and fallback procedures
- Anomaly detection systems

- Human override capabilities
- Emergency shutdown procedures
- Safety testing results

■ **Explainability Tools**

- Description of explanation methods used
- Examples of explanations generated
- User testing results for explanation clarity
- Technical documentation for auditors
- API documentation if explanations available programmatically

E.2 Bias & Demographic Fairness

■ Pre-Deployment Bias Testing

- Statistical parity analysis across demographic groups
- Equalized odds testing
- Calibration assessment by group
- Intersectional bias testing
- Documentation of protected characteristics tested

■ Proxy Variable Identification

- List of potential proxy variables identified
- Correlation analysis with protected characteristics
- Testing of model with proxy variables perturbed
- Documentation of proxy-aware fairness constraints

■ Regional Lexicon Integration

- Inventory of language variants included in training
- Demographic representativeness of training data
- Validation dataset demographic composition
- Testing results on diverse populations

■ Continuous Monitoring

- Bias monitoring system implementation
- Automated alerting for bias detection
- Model drift detection mechanisms
- Feedback loop analysis
- Schedule for periodic reauditing

■ Disaggregated Performance Metrics (Tier 3+)

- Performance metrics by demographic group
- Intersectional performance analysis
- Gap analysis between groups
- Minimum performance thresholds defined and met

■ Mitigation Strategies

- Documentation of bias mitigation approaches
- Effectiveness evidence for mitigations

- Residual bias quantification
- Ongoing mitigation plans

E.3 Environmental Impact Assessment

■ Training Phase Carbon Footprint

- Total compute hours used in training
- Hardware specifications (GPUs, TPUs, etc.)
- Energy consumption calculation (kWh)
- Grid carbon intensity for training location
- Total CO₂ equivalent emissions from training

■ Operational Phase Assessment

- Per-inference energy consumption
- Estimated total operational emissions
- Data center PUE (Power Usage Effectiveness)
- Expected usage volume and scaling projections

■ Hardware Lifecycle

- Embodied carbon in specialized hardware
- Expected hardware lifetime
- E-waste management and disposal plans

■ Efficiency Measures Implemented

- Model architecture efficiency optimizations
- Training efficiency techniques used
- Inference optimization (compression, caching, etc.)
- Infrastructure efficiency choices
- Renewable energy usage

■ Public Disclosure

- Carbon footprint disclosed publicly
- Methodology for calculations transparent
- Updated annually or more frequently

E.4 Neuro-Rights Compliance (if applicable)

■ Consent Mechanism

- Informed consent documentation prepared
- Consent covers all aspects of neural data collection
- Real-time revocation mechanism implemented
- Consent testing with users completed

■ Actionable Input Detection

- Emotional/sensory detection system implemented
- Detection accuracy validated
- Human review protocol established
- Reviewer training completed
- 24/7 reviewer availability ensured

■ Data Protection

- Enhanced encryption implemented
- Access controls stricter than standard data
- Anonymization where feasible
- Breach response plan specific to neural data

■ Neuro-Ethics Audit

- Quarterly audit schedule established
- Neuro-ethics committee constituted or external auditor engaged
- Initial audit completed with findings addressed
- Audit methodology documented

E.5 Whistleblower System

■ Anonymous Portal

- Encrypted reporting portal implemented
- End-to-end encryption verified
- Anonymity preservation tested
- Multi-language support confirmed
- Accessibility compliance verified

■ 72-Hour Escalation

- Escalation process documented
- Timekeeping and alerting system implemented
- External oversight body identified and engaged
- Automatic notification system on failure tested
- Staff trained on escalation requirements

■ Protection Measures

- 120-day role protection policy documented
- Anti-retaliation enforcement authority established
- Protection verification process implemented
- Staff and management trained on protections

■ Testing

- Mock reporting scenarios conducted
- Portal penetration tested by independent security auditors
- Escalation process tested end-to-end
- Response times verified

E.6 Immutable Audit Trails

■ Technical Implementation

- Cryptographic hashing implemented
- Write-once storage configured
- Distributed verification capability
- ISO/IEC 27001 or equivalent compliance verified

■ Content Completeness

- All required events logged (decisions, changes, access, incidents)
- Log format structured and parseable
- Sufficient detail for reconstruction of system behavior
- Timestamp accuracy verified

■ Access Controls

- Role-based access to logs implemented
- Access auditing enabled
- Regulator access procedures documented
- Rapid extraction capability verified

■ Retention Compliance

- Retention periods configured per tier requirements
- Backup and disaster recovery for logs
- Legal hold processes documented
- Log deletion only after retention period (if no legal hold)

E.7 Ombudsman Process (Tier 2+)

■ Ombudsman Access

- Ombudsman contact information publicized
- Multiple submission channels available
- No-cost access verified
- Multi-language support confirmed

■ Process Documentation

- Challenge process clearly documented
- Timelines for each stage defined
- Expected outcomes explained
- User guidance materials prepared

■ Organizational Cooperation

- Internal procedures for ombudsman requests established
- Document production timelines confirmed
- Staff trained on cooperation requirements
- Remediation implementation processes ready

E.8 Cross-Certification (if applicable)

■ Framework Identification

- All applicable frameworks identified (GDPR, CPRA, LGPD, etc.)
- Jurisdictional requirements mapped
- Gaps between GAAP and frameworks documented

■ Supplementary Compliance

- Framework-specific requirements not covered by GAAP identified
- Supplementary audits scheduled or completed
- Cross-recognition agreements verified

■ Documentation Alignment

- Single documentation set covering all frameworks where possible
- Framework-specific documentation prepared where needed
- Version control across frameworks maintained

8. Glossary & Definitions

Actionable Inputs: Neural signals from brain-computer interfaces indicating emotional states, distress, or sensory experiences that trigger mandatory human review rather than automated AI response. Systems must pause automated decision-making and notify human operators when actionable inputs are detected.

Bias Mitigation: Techniques and processes used to identify, measure, and reduce unfair discrimination in AI systems across demographic groups. Includes pre-processing (improving training data balance), in-processing (fairness constraints during model training), and post-processing (output calibration) approaches.

Brain-Computer Interface (BCI): Technology enabling direct communication between the brain and external devices, typically through recording neural signals and translating them into commands or data. Includes both invasive (implanted) and non-invasive (surface) interfaces.

Cognitive Liberty: The right to mental self-determination, including freedom of thought, mental privacy, and protection from non-consensual cognitive manipulation or monitoring. A fundamental neuro-right in the context of AI and BCI systems.

Cross-Border Data Transfer: Movement of personal data or AI training data from one legal jurisdiction to another, requiring specific legal bases, transparency, and often explicit consent due to varying data protection regimes.

Data Localization: Legal requirement that certain categories of data must be stored and processed within specific geographic or jurisdictional boundaries, often for national security, privacy, or economic policy reasons.

Data Sovereignty: The principle that data is subject to the laws and governance structures of the jurisdiction where it is collected or where the data subject resides. In indigenous context, refers to collective rights of indigenous peoples to control data about their communities.

Demographic Fairness: Achievement of equitable outcomes and performance across different demographic groups in AI system operation. Requires disaggregated testing and validation to ensure no group experiences systematically worse outcomes.

Disaggregated Metrics: Performance measurements reported separately for different demographic groups rather than only averaged across the entire population. Essential for detecting bias that may be hidden in aggregate statistics.

Embodied Carbon: Total greenhouse gas emissions generated during the manufacturing, transportation, and disposal of physical goods (in AI context: specialized hardware like GPUs, TPUs, data center equipment). Distinct from operational carbon from energy consumption during use.

Explainability: The degree to which AI system decision-making processes can be understood and described by humans. Includes both technical explainability (for experts/auditors) and user-facing explainability (for affected individuals).

Free, Prior, and Informed Consent (FPIC): A principle in international human rights law requiring meaningful consent from indigenous peoples before undertaking projects or activities that may affect them, their territories, or their rights. Must be obtained through culturally appropriate processes respecting indigenous governance structures.

Immutable Audit Logs: Tamper-proof records of AI system behavior and operations that cannot be altered or deleted after creation. Typically implemented using cryptographic techniques and write-once storage, adhering to standards like ISO/IEC 27001.

Intersectional Analysis: Examination of how multiple aspects of identity (race, gender, age, disability, etc.) combine to create unique experiences of discrimination or disadvantage. Essential for comprehensive bias testing as combinations of characteristics may reveal issues not visible when examining single dimensions.

Machine Unlearning: Techniques for removing the influence of specific training data from a trained machine learning model without complete retraining. Important for honoring data deletion requests and protecting privacy.

Model Drift: Changes in AI system behavior over time, often due to shifts in input data distribution, changes in the real-world environment, or feedback loops. Requires continuous monitoring to detect and address.

Neural Lag: Time delay between brain activity and AI system response in brain-computer interface applications. Excessive lag can cause loss of user control, cognitive dissonance, or unsafe operation.

Neuro-Rights: Fundamental rights related to neural data, cognitive liberty, and brain integrity. Include mental privacy, freedom from cognitive manipulation, and protection of personal neural information.

OCAP® Principles: Framework for indigenous data governance standing for Ownership, Control, Access, and Possession. Asserts that indigenous peoples own their collective data and have rights to control all aspects of data collection, management, and use affecting their communities.

Ombudsman: Independent authority empowered to investigate complaints, mediate disputes, and issue binding decisions regarding AI system impacts on individuals. Provides recourse mechanism for affected persons without requiring legal representation.

Proxy Variable: Data feature that does not explicitly represent a protected characteristic (race, gender, etc.) but correlates with it and can thus serve as an indirect pathway for discrimination. Examples include ZIP code (race, socioeconomic status), name patterns (ethnicity, gender), or educational institution (socioeconomic background).

Rebuttable Presumption: Legal principle where a certain fact is assumed to be true unless the opposing party presents evidence to disprove it. In GAAP context, adverse employment actions against whistleblowers within 120 days are presumed to be retaliation unless the organization proves otherwise.

Responsible AI Officer (RAIO): Designated organizational role with independent authority to oversee AI ethics and compliance. Required for Tier 2+ systems. Must have genuine power to block deployments, access to leadership, and protection from retaliation.

Secular (in GAAP context): Universal applicability regardless of cultural, religious, or philosophical backgrounds. GAAP's secular approach means it establishes ethical requirements based on empirically demonstrable principles rather than specific ethical traditions, enabling global adoption without exclusion.

Tier (GAAP Tiering System): Classification system scaling requirements to AI system risk and impact. Tier 1 (local, low impact) through Tier 4 (military, governmental, critical infrastructure) with increasingly stringent requirements for higher-risk systems.

Whistleblower: Individual who reports violations, safety concerns, or ethical issues regarding AI systems, typically from within the organization developing or deploying them. GAAP 2.2 provides extensive protections including anonymity, mandatory escalation, and anti-retaliation measures.

9. Contact & Oversight

9.1 International Ethics Accountability Body (IEAB)

The IEAB serves as the central coordinating authority for GAAP implementation globally.

Primary Functions:

- Accreditation of national and regional certification bodies
- Maintenance of unified audit methodology and standards
- Cross-recognition agreement facilitation
- Dispute resolution and interpretation guidance
- Coordination with regulatory authorities
- Quality assurance for certification processes

Contact Information:

- Website: www.ieab-gaap.org (illustrative - actual implementation would require establishment)
- Email: certification@ieab-gaap.org
- Phone: International hotline (details to be established)
- Physical Address: [To be determined - likely Geneva, Switzerland for neutrality]

Reporting Violations:

- Anonymous violation reporting portal: report.ieab-gaap.org
- Encrypted communication via PGP: [Public key to be published]
- Whistleblower protection guaranteed for good-faith reports

9.2 Regional Certification Authorities

IEAB-accredited bodies in major jurisdictions:

North America:

- United States: AI Safety Commission (ASC)
- Canada: Office of the AI Ombudsman
- Mexico: Autoridad de Inteligencia Artificial Responsable

Europe:

- European Union: EU AI Ethics Board (coordinates with EU AI Act authorities)

- United Kingdom: Centre for Data Ethics and Innovation (CDEI)
- Switzerland: Federal AI Accountability Office

Asia-Pacific:

- Singapore: AI Governance Committee
- Australia: Responsible AI Framework Authority
- Japan: AI Principles Implementation Office
- India: National AI Ethics Board

Other Regions:

- Brazil: Autoridade Nacional de Proteção de Dados (ANPD) - AI Division
- South Africa: Information Regulator AI Unit
- Middle East: Regional AI Governance Council (multi-country)

Contact details: Available at www.ieab-gaap.org/regional-authorities

9.3 Enforcement Contacts

For immediate safety concerns or urgent violations:

Emergency Hotline: [To be established - 24/7 multilingual support]

- Immediate response for Tier 4 system failures
- Rapid deployment of emergency auditors
- Authority to order immediate system suspension

Regulatory Inquiries:

- Policy interpretation requests: policy@ieab-gaap.org
- Certification status verification: verify@ieab-gaap.org
- Audit scheduling and coordination: audits@ieab-gaap.org

9.4 Public Resources

Documentation and Guidance:

- Full GAAP 2.2 specification: www.ieab-gaap.org/standard
- Implementation guides by tier: www.ieab-gaap.org/guides
- Training materials for RAIOS and auditors: www.ieab-gaap.org/training
- Case studies and best practices: www.ieab-gaap.org/case-studies

Transparency Database:

- Certified systems registry: www.ieab-gaap.org/registry
- Violation reports and resolutions: www.ieab-gaap.org/violations
- Ombudsman decision database: www.ieab-gaap.org/ombudsman
- Statistical dashboards: www.ieab-gaap.org/statistics

9.5 Stakeholder Engagement

Industry Forum:

- Quarterly meetings for certified organizations
- Sharing of implementation challenges and solutions
- Input into standard evolution
- Networking and collaboration opportunities
- Register: forum@ieab-gaap.org

Civil Society Consultation:

- Open consultation periods for standard updates
- Representation in governance bodies
- Advocacy and watchdog functions
- Public comment opportunities
- Engage: civilsociety@ieab-gaap.org

Academic Collaboration:

- Research partnerships
- Validation of methodologies
- Development of technical standards
- Education and training programs
- Collaborate: research@ieab-gaap.org

9.6 Feedback and Standard Evolution

GAAP is a living standard that evolves with technology and social understanding:

Feedback Channels:

- General feedback: feedback@ieab-gaap.org
- Technical suggestions: technical@ieab-gaap.org

- Proposed amendments: Submit via standard revision process documented at www.ieab-gaap.org/amendments

Version Updates:

- Minor updates (X.X): Annually, incorporating clarifications and small improvements
- Major updates (X.0): Every 2-3 years, incorporating significant new requirements or structural changes
- All updates subject to public comment period
- Backward compatibility and transition periods provided

Notification:

Subscribe to updates: subscribe@ieab-gaap.org

- Email notifications of new versions
- Webinars on significant changes
- Transition guidance for maintaining compliance

10. References

This white paper draws on extensive research, international standards, and ethical frameworks:

AI Ethics Frameworks:

- IEEE, "Ethically Aligned Design," IEEE Standards Association, 2019
- European Commission, "Ethics Guidelines for Trustworthy AI," 2019
- OECD, "Recommendation of the Council on Artificial Intelligence," 2019
- Montreal Declaration for Responsible AI, 2018

Data Protection and Privacy:

- General Data Protection Regulation (GDPR), EU 2016/679
- California Privacy Rights Act (CPRA), 2020
- Lei Geral de Proteção de Dados (LGPD), Brazil, 2018
- Personal Information Protection and Electronic Documents Act (PIPEDA), Canada

Technical Standards:

- ISO/IEC 23894:2023, Information technology — Artificial intelligence — Risk management
- ISO/IEC 27001:2022, Information security management systems
- IEEE 7000-2021, Model Process for Addressing Ethical Concerns During System Design
- IEEE 7010-2020, Recommended Practice for Assessing the Impact of Autonomous and Intelligent Systems on Human Well-Being

Fairness and Bias:

- Barocas, S., Hardt, M., & Narayanan, A., "Fairness and Machine Learning: Limitations and Opportunities," 2019
- Mehrabi, N., et al., "A Survey on Bias and Fairness in Machine Learning," ACM Computing Surveys, 2021
- Buolamwini, J. & Gebru, T., "Gender Shades: Intersectional Accuracy Disparities in Commercial Gender Classification," 2018

Explainability and Transparency:

- Ribeiro, M.T., et al., "Why Should I Trust You?: Explaining the Predictions of Any Classifier," KDD 2016
- Lundberg, S. & Lee, S., "A Unified Approach to Interpreting Model Predictions," NeurIPS 2017
- Lipton, Z., "The Mythos of Model Interpretability," Communications of the ACM, 2018

Neuro-Rights and BCI Ethics:

- Yuste, R., et al., "Four Ethical Priorities for Neurotechnologies and AI," Nature, 2017
- Ienca, M. & Andorno, R., "Towards New Human Rights in the Age of Neuroscience and Neurotechnology," Life Sciences, Society and Policy, 2017
- Goering, S. & Klein, E., "Are Neurorights Human Rights?", AJOB Neuroscience, 2023

Indigenous Data Sovereignty:

- First Nations Information Governance Centre, "OCAP®: Ownership, Control, Access and Possession," 2022
- Carroll, S.R., et al., "The CARE Principles for Indigenous Data Governance," Data Science Journal, 2020
- Research Data Alliance International Indigenous Data Sovereignty Interest Group, "CARE Principles," 2019

Environmental Sustainability:

- Strubell, E., et al., "Energy and Policy Considerations for Deep Learning in NLP," ACL 2019
- Patterson, D., et al., "Carbon Emissions and Large Neural Network Training," arXiv, 2021
- Schwartz, R., et al., "Green AI," Communications of the ACM, 2020

Whistleblower Protections:

- Government Accountability Project, "Best Practices for Whistleblower Policies," 2018
- OECD, "Committing to Effective Whistleblower Protection," 2016
- Transparency International, "International Principles for Whistleblower Legislation," 2013

Governance and Accountability:

- Floridi, L., et al., "AI4People—An Ethical Framework for a Good AI Society," Minds and Machines, 2018
- Bryson, J. & Winfield, A., "Standardizing Ethical Design for Artificial Intelligence and Autonomous Systems," Computer, 2017
- Raji, I.D., et al., "Closing the AI Accountability Gap," FAT* 2020

Acknowledgments

GAAP 2.2 was developed through extensive multi-stakeholder collaboration including:

- AI researchers and developers across academia and industry
- Ethicists specializing in technology, data, and human rights
- Policymakers and regulators from multiple jurisdictions
- Civil society organizations and advocacy groups
- Affected communities and user representatives
- Legal experts in privacy, human rights, and technology law
- Indigenous rights advocates and community leaders
- Environmental sustainability experts
- Neuroscientists and neuro-ethics specialists

Special recognition to organizations that contributed to stress testing, validation, and refinement of this framework.

Version History

GAAP 1.0 (2022): Initial framework establishing core transparency and fairness principles

GAAP 2.0 (2024): Introduction of tiered implementation structure and enhanced bias mitigation

GAAP 2.1 (Early 2025): Addition of preliminary environmental considerations and cross-border governance

GAAP 2.2 (July 2025): Current version incorporating:

- Comprehensive neuro-rights and BCI framework
- Enhanced whistleblower protections with mandatory 72-hour escalation
- Contextual harm assessment based on regional support availability
- Indigenous data sovereignty requirements
- Global interadoptability with major regulatory frameworks
- Strengthened enforcement mechanisms

Future Development: GAAP 3.0 anticipated 2027-2028 to address emerging challenges in artificial general intelligence, quantum AI, and evolving governance needs.

License and Usage

This document (GAAP 2.2 White Paper) is released under Creative Commons Attribution 4.0 International (CC BY 4.0).

You are free to:

- Share: Copy and redistribute the material in any medium or format
- Adapt: Remix, transform, and build upon the material for any purpose

Under the following terms:

- Attribution: You must give appropriate credit to IEAB and GAAP, provide a link to the license, and indicate if changes were made

The GAAP standard itself is maintained by the International Ethics Accountability Body (IEAB) and may be adopted, adapted, or referenced by organizations worldwide.

For questions about this white paper or GAAP 2.2 implementation, please contact the International Ethics Accountability Body at info@ieab-gaap.org

Document Version: 1.0

Last Updated: July 2025

Next Review: July 2026